

# Terraform Walk-Through

# Our Plan

- ◆ Getting started with Terraform
  - Setting up your AWS account
  - Installing Terraform
  - Deploying a single server
  - Deploying a single web server
- ◆ More advanced configuration
  - Deploying a configurable web server
  - Deploying a cluster of web servers
  - Deploying a load balancer
  - Cleaning up

# We Will Use AWS

- ◆ The largest market share
  - Although for you it may be different
- ◆ Will try to provide examples for other clouds
- ◆ All clouds give you a free tier
  - A trick to user after a year:
  - Use kind of address
  - If you already used up your free tier credits, the examples in the labs should still cost you no more than a few dollars.

# Setting Up Your AWS Account

- ◆ If you don't already have an **AWS account**, head over to <https://aws.amazon.com> and sign up
- ◆ The *only* thing you should use the root user for is to create other user accounts with more-limited permissions, and then switch to one of those accounts immediately
- ◆ If you are using an existing AWS account, it must have a **Default VPC** in it.
- ◆ If the instructor provided a student account, you can use that

Add user

## Set user details

You can add multiple users at once with the same access type and permissions. [Learn more](#)

User name\*

[+](#) Add another user

## Select AWS access type

Select how these users will access AWS. Access keys and autogenerated passwords are provided in the last step

Access type\*



**Programmatic access**

Enables an **access key ID** and **secret access key** for the AWS API, other development tools.



**AWS Management Console access**

Enables a **password** that allows users to sign-in to the AWS Manage

Console password\*



Autogenerated password



Custom password

Require password reset




User must create a new password at next sign-in


# Giving AWS Permissions


- ◆ Here are the permissions you will need (for some labs)
  - AmazonEC2FullAccess
  - AmazonS3FullAccess
  - AmazonDynamoDBFullAccess
  - AmazonRDSFullAccess
  - CloudWatchFullAccess
  - IAMFullAccess

Add user 1 2 3 4 5

▼ Set permissions

 Add user to group

 Copy permissions from existing user

 Attach existing policies directly

Add user to an existing group or create a new one. Using groups is a best-practice way to manage user's permissions by job functions. [Learn more](#)

Add user to group

Create group

Refresh

Search

Showing 4 results

Group ▼	Attached policies
<input checked="" type="checkbox"/> admins	AdministratorAccess
<input type="checkbox"/> devs	AmazonRDSFullAccess and 2 more
<input type="checkbox"/> external_trainers	SecurityGroupsAccess
<input type="checkbox"/> in_class	AmazonEC2FullAccess and 1 more

► Set permissions boundary

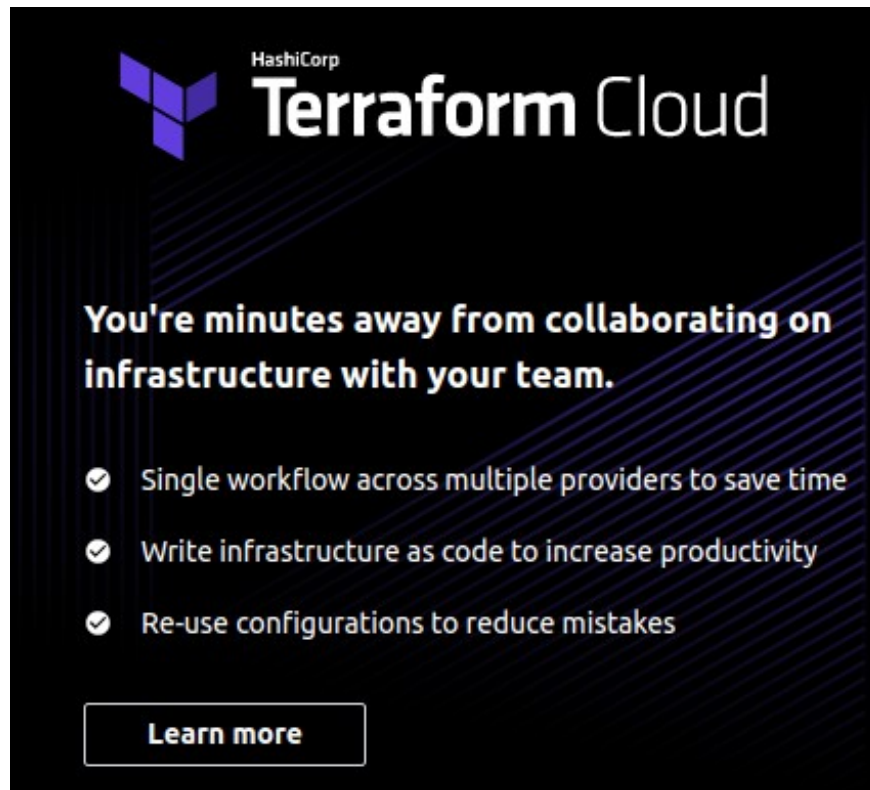
Cancel

Previous

Next: Tags

# Install Terraform

- ◆ Download the terraform executable from the [Terraform home page](#)
- ◆ Depending on your OS, you may also install a native package
- ◆ For Mac
  - `brew install terraform`
- ◆ You may use a cloud server if provided by the instructor



# Verify

- ◆ Ubuntu example
- ◆ Terraform v0.12.20
- ◆ Your version of Terraform is out of date! The latest version is 0.12.24. You can update by downloading from <https://www.terraform.io/downloads.html>
- ◆ OK... Update [here](#)

# Now What?

```
(base) mark@mark-workstation:~$ terraform
Usage: terraform [-version] [-help] <command> [args]
```

The available commands for execution are listed below.  
The most common, useful commands are shown first, followed by less common or more advanced commands. If you're just getting started with Terraform, stick with the common commands. For the other commands, please read the help and docs before usage.

Common commands:

apply	Builds or changes infrastructure
console	Interactive console for Terraform interpolations
destroy	Destroy Terraform-managed infrastructure
env	Workspace management
fmt	Rewrites config files to canonical format
get	Download and install modules for the configuration
graph	Create a visual graph of Terraform resources
import	Import existing infrastructure into Terraform
init	Initialize a Terraform working directory
login	Obtain and save credentials for a remote host
logout	Remove locally-stored credentials for a remote host
output	Read an output from a state file
plan	Generate and show an execution plan
providers	Prints a tree of the providers used in the configuration
refresh	Update local state file against real resources
show	Inspect Terraform state or plan
taint	Manually mark a resource for recreation
untaint	Manually unmark a resource as tainted
validate	Validates the Terraform files
version	Prints the Terraform version
workspace	Workspace management



# Connect to AWS - Way 1

- ◆ Way 1: Set up AWS credentials
- ◆ `$export AWS_ACCESS_KEY_ID=(you access key id)`
- ◆ `$export AWS_SECRET_ACCESS_KEY=(your secret access key)`
- ◆ This will only give you the setup for this shell
- ◆ To make it work after reboot put it into `.bashrc`
- ◆ `vi ~/.bashrc`

# Connect to AWS - Way 2

- ◆ Way 2
- ◆ Terraform supports the same authentication mechanisms as all AWS CLI and SDK tools
- ◆ Therefore, it'll also be able to use credentials in `$HOME/.aws/credentials`
  - These are automatically generated if you run the `configure` command on the AWS CLI, or IAM

# Let Us Prepare to Deploy a Server

- ◆ Create an empty folder, lab01
- ◆ Put in a file called main.tf
- ◆ Now you will be adding resources, like this:

```
resource "<PROVIDER>_<TYPE>" "<NAME>" {  
    [CONFIG ...]  
}
```

# Server Resource

- ◆ ami
  - The Amazon Machine Image (AMI) to run on the EC2 Instance.
  - You can find free and paid AMIs in the AWS Marketplace
  - or create your own using tools such as Packer
  - This ami parameter to the ID of an Ubuntu 18.04 AMI in us-east-1  
This AMI is free to use
- ◆ instance\_type
  - The type of EC2 Instance to run
  - Each type of EC2 Instance provides a different amount of CPU, memory, disk space, and networking capacity.
  - The EC2 Instance Types page lists all the available options
  - t2.micro, which has one virtual CPU, 1 GB of memory, and is part of the AWS free tier

# Do Not Try to Remember by Heart

- ◆ Terraform supports dozens of providers
- ◆ Each of which supports dozens of resources
- ◆ Each resource has dozens of arguments
- ◆ We recommend using the documentation.
  - Here is an [example for ami](#)
- ◆ Now run `terraform init`

# Result of

```
(base) mark@mark-workstation:~/projects/ES/terraform-class/lab01$ terraform init
```

Initializing the backend...

Initializing provider plugins...

- Checking for available provider plugins...
- Downloading plugin for provider "aws" (hashicorp/aws) 2.60.0...

The following providers do not have any version constraints in configuration, so the latest version was installed.

To prevent automatic upgrades to new major versions that may contain breaking changes, it is recommended to add version = "..." constraints to the corresponding provider blocks in configuration, with the constraint strings suggested below.

```
* provider.aws: version = "~> 2.60"
```

**Terraform has been successfully initialized!**

You may now begin working with Terraform. Try running "terraform plan" to see any changes that are required for your infrastructure. All Terraform commands should now work.

If you ever set or change modules or backend configuration for Terraform, rerun this command to reinitialize your working directory. If you forget, other commands will detect it and remind you to do so if necessary.

# Result of "terraform plan"

```
1 $ terraform plan Refreshing
2 Terraform state in-memory prior to plan...
3 The refreshed state will be used to calculate this plan, but will not be
4 persisted to local or remote state storage.
5 -----
6 An execution plan has been generated and is shown below.
7 Resource actions are indicated with the following symbols:
8   + create
9
10 Terraform will perform the following actions:
11
12   # aws_instance.example will be created
13   + resource "aws_instance" "example" {
14       + ami                               = "ami-0c55b159cbfaffe1f0"
15 Plan: 1 to add, 0 to change, 0 to destroy.
16
17 -----
18
19 Note: You didn't specify an "-out" parameter to save this plan, so Terraform
20 can't guarantee that exactly these actions will be performed if
21 "terraform apply" is subsequently run.
```

# Result of "terraform apply"

```
1 An execution plan has been generated and is shown below.
2 Resource actions are indicated with the following symbols:
3   + create
4
5 Terraform will perform the following actions:
6
7   # aws_instance.example will be created
8   + resource "aws_instance" "example" {
9       + ami                                = "ami-0c55b159cbfaffe1f0"
10  ...
11  }
12
13 Plan: 1 to add, 0 to change, 0 to destroy.
14
15 Do you want to perform these actions?
16   Terraform will perform the actions described above.
17   Only 'yes' will be accepted to approve.
18
19   Enter a value: yes
20
21 aws_instance.example: Creating...
22 aws_instance.example: Still creating... [10s elapsed]
23 aws_instance.example: Still creating... [20s elapsed]
24 aws_instance.example: Creation complete after 24s [id=i-0a4dfc4992739cb6e]
25
26 Apply complete! Resources: 1 added, 0 changed, 0 destroyed.
```



# Verify the Deployment Result

- ◆ Go to AWS dashboard
- ◆ Verify that the server was indeed created

Launch Instance

▼

Connect

Actions ▼

🔍

Filter by tags and attributes or search by keyword

<input type="checkbox"/>	Name ▲	Instance ID ▼	Instance Type ▼	Availability Zone ▼	Instance State ▼	Status Checks ▼
<input type="checkbox"/>		i-0a4dfc4992739cb6e	t2.micro	us-east-2c	<span>●</span> running	<span>✓</span> 2/2 checks ...

# Next Step

- ◆ Let's give our server a name tag
- ◆ Add to your `main.tf`
  - (remove the previous server definition)

```
1 resource "aws_instance" "example" {
2     ami           = "ami-0c55b159cbfafa1f0"
3     instance_type = "t2.micro"
4
5     tags = {
6         Name = "terraform-example"
7     }
8 }
```

- ◆ Run `terraform apply`
- ◆ Verify that we gave our server a name

```
Plan: 2 to add, 0 to change, 1 to destroy.

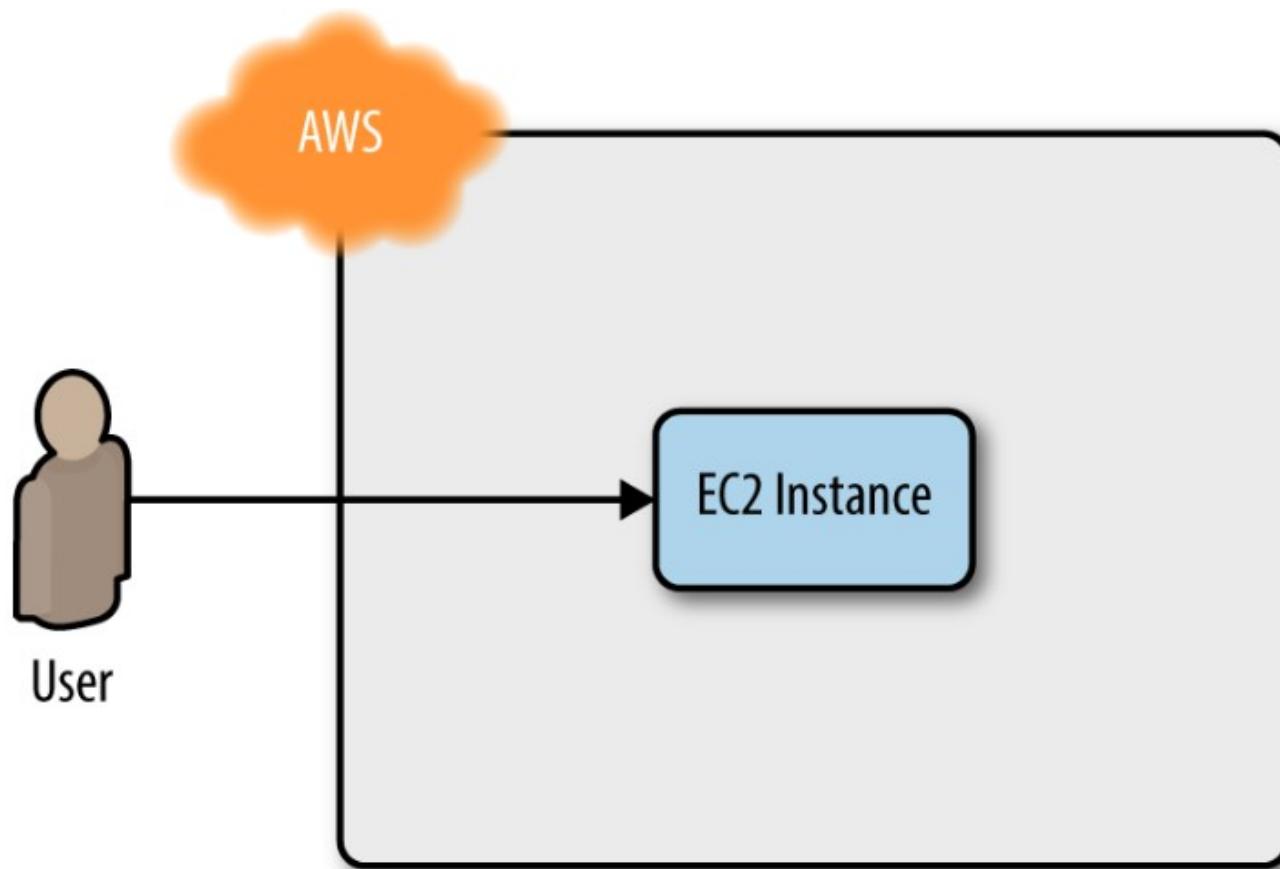
Do you want to perform these actions?
  Terraform will perform the actions described above.
  Only 'yes' will be accepted to approve.

  Enter a value: yes

aws_instance.example: Destroying... [id=i-0a4dfc4992739cb6e]
aws_security_group.instance: Creating...
aws_security_group.instance: Creation complete after 2s [id=sg-02b2ec7c6f419809a]
aws_instance.example: Still destroying... [id=i-0a4dfc4992739cb6e, 10s elapsed]
aws_instance.example: Still destroying... [id=i-0a4dfc4992739cb6e, 20s elapsed]
aws_instance.example: Still destroying... [id=i-0a4dfc4992739cb6e, 30s elapsed]
aws_instance.example: Destruction complete after 31s
aws_instance.example: Creating...
aws_instance.example: Still creating... [10s elapsed]
aws_instance.example: Still creating... [20s elapsed]
aws_instance.example: Creation complete after 23s [id=i-01f311884533ef669]

Apply complete! Resources: 2 added, 0 changed, 1 destroyed.
```

# Deploy a Single Web Server



# Make a Web Server

- ◆ In the real world, you would build a real web server such as Flask
- ◆ We will, instead, do a one-command

```
1 #!/bin/bash
2 echo "Hello, World" > index.html
3 nohup busybox httpd -f -p 8080 &
```

- ◆ However, how should we put it into our instance?
- ◆ We will add it to the `aws_instance`, as *User Data* configuration

# Adding a Script to the Instance

- ◆ You pass a shell script to User Data by setting the `user_data` argument in your Terraform code as follows:
- ◆ The `<<-EOF` and `EOF` are Terraform's heredoc syntax, which allows you to create multiline strings without having to insert newline characters all over the place

```
1 user_data = <<-EOF
2             #!/bin/bash
3             echo "Hello, World" > index.html
4             nohup busybox httpd -f -p 8080 &
5             EOF
```

# Wait! One More Thing!

- ◆ By default, AWS does not allow any incoming or outgoing traffic from an EC2 Instance.
- ◆ To allow the EC2 Instance to receive traffic on port 8080, you need to create a security group:
- ◆ Creates a new resource called `aws_security_group`

```
1 resource "aws_security_group" "instance" {  
2   name = "terraform-example-instance"  
3  
4   ingress {  
5     from_port = 8080  
6     to_port   = 8080  
7     protocol  = "tcp"  
8     cidr_blocks = ["0.0.0.0/0"]  
9   }  
10 }
```

# CIDR Blocks

- ◆ The `ingress` in this group allows incoming TCP requests
  - on port 8080 from the CIDR block `0.0.0.0/0`
- ◆ CIDR blocks are a concise way to specify IP address ranges
- ◆ For example
  - a CIDR block of `10.0.0.0/24`
  - represents all IP addresses between `10.0.0.0` and `10.0.0.255`
- ◆ The CIDR block `0.0.0.0/0` is an IP address range that includes all possible IP addresses, so this security group allows incoming requests on port 8080 from any IP

# Passing the Security Group ID

- ◆ You also need to tell the EC2 instance to actually use the security group by passing the group's ID
- ◆ It goes into the `vpc_security_group_ids` argument of the `aws_instance` resource.
- ◆ This is done with Terraform expressions



# Terraform "expressions"

- ◆ An expression in Terraform is anything that returns a value
- ◆ The simplest type of expressions are literals
  - strings: "ami-0c55b159cbfafa1f0"
  - numbers: 7
- ◆ Here we need an expression which is a reference

```
1 <PROVIDER>_<TYPE>.<NAME>.<ATTRIBUTE>  
2  
3 In our case  
4  
5 aws_security_group.instance.id
```

# Altogether

```
resource "aws_instance" "example" {
  ami                = "ami-0c55b159cbfafa1f0"
  instance_type      = "t2.micro"
  vpc_security_group_ids = [aws_security_group.instance.id]

  user_data = <<-EOF
    #!/bin/bash
    echo "Hello, World" > index.html
    nohup busybox httpd -f -p 8080 &
  EOF

  tags = {
    Name = "terraform-example"
  }
}
```

# New Result of "terraform plan"

```
1 Refreshing Terraform state in-memory prior to plan...
2 The refreshed state will be used to calculate this plan, but will not be
3 persisted to local or remote state storage.
4
5 aws_instance.example: Refreshing state... [id=i-0a4dfc4992739cb6e]
6 ...
7   # aws_security_group.instance will be created
8   + resource "aws_security_group" "instance" {
9     ...
10 Plan: 2 to add, 0 to change, 1 to destroy.
```

# New Result of "terraform apply"

```
Plan: 2 to add, 0 to change, 1 to destroy.

Do you want to perform these actions?
  Terraform will perform the actions described above.
  Only 'yes' will be accepted to approve.

  Enter a value: yes

aws_instance.example: Destroying... [id=i-0a4dfc4992739cb6e]
aws_security_group.instance: Creating...
aws_security_group.instance: Creation complete after 2s [id=sg-02b2ec7c6f419809a]
aws_instance.example: Still destroying... [id=i-0a4dfc4992739cb6e, 10s elapsed]
aws_instance.example: Still destroying... [id=i-0a4dfc4992739cb6e, 20s elapsed]
aws_instance.example: Still destroying... [id=i-0a4dfc4992739cb6e, 30s elapsed]
aws_instance.example: Destruction complete after 31s
aws_instance.example: Creating...
aws_instance.example: Still creating... [10s elapsed]
aws_instance.example: Still creating... [20s elapsed]
aws_instance.example: Creation complete after 23s [id=i-01f311884533ef669]

Apply complete! Resources: 2 added, 0 changed, 1 destroyed.
```

# Et Voila!

Launch Instance ▼

Connect

Actions ▼

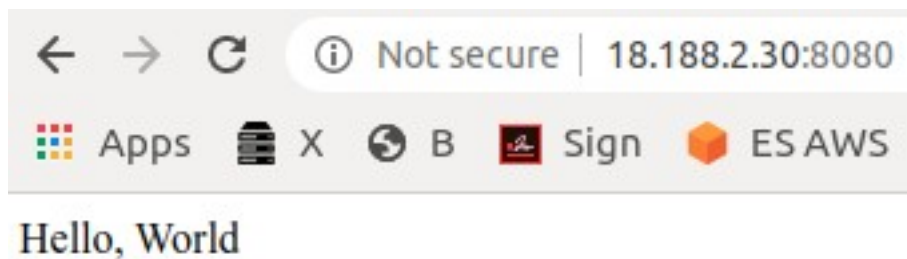
🔍 Filter by tags and attributes or search by keyword

<input type="checkbox"/>	Name ▲	Instance ID ▼	Instance Type ▼	Availability Zone ▼	Instance State ▼
<input type="checkbox"/>	terraform-example	i-01f311884533ef669	t2.micro	us-east-2c	● running
<input type="checkbox"/>	terraform-example	i-0a4dfc4992739cb6e	t2.micro	us-east-2c	● terminated
<input type="checkbox"/>		i-0f7e5b9e4508a28cb	t2.micro	us-east-2b	● running

# Test the Deployment

```
1 $ curl http://<EC2_INSTANCE_PUBLIC_IP>:8080
2 Hello, World
```

```
1 $ curl http://18.188.2.30:8080
2 Hello, World
```



# Quiz

- ◆ Usernames and passwords referenced in the Terraform code, even as variables, will end up in plain text in the state file.
  - A. True
  - B. False

# Quiz

- ◆ What happens when you apply Terraform configuration?  
Choose *TWO* correct answers.
  - A. terraform plan
  - B. terraform state
  - C. terraform apply
  - D. terraform validate
  - E. terraform output



# Terraform Dependencies

- ◆ When you add a reference from one resource to another, you create an implicit dependency
- ◆ Terraform
  - Parses these dependencies
  - builds a dependency graph from them
  - uses that to automatically determine in which order it should create resources
- ◆ To see the dependencies, you use the command

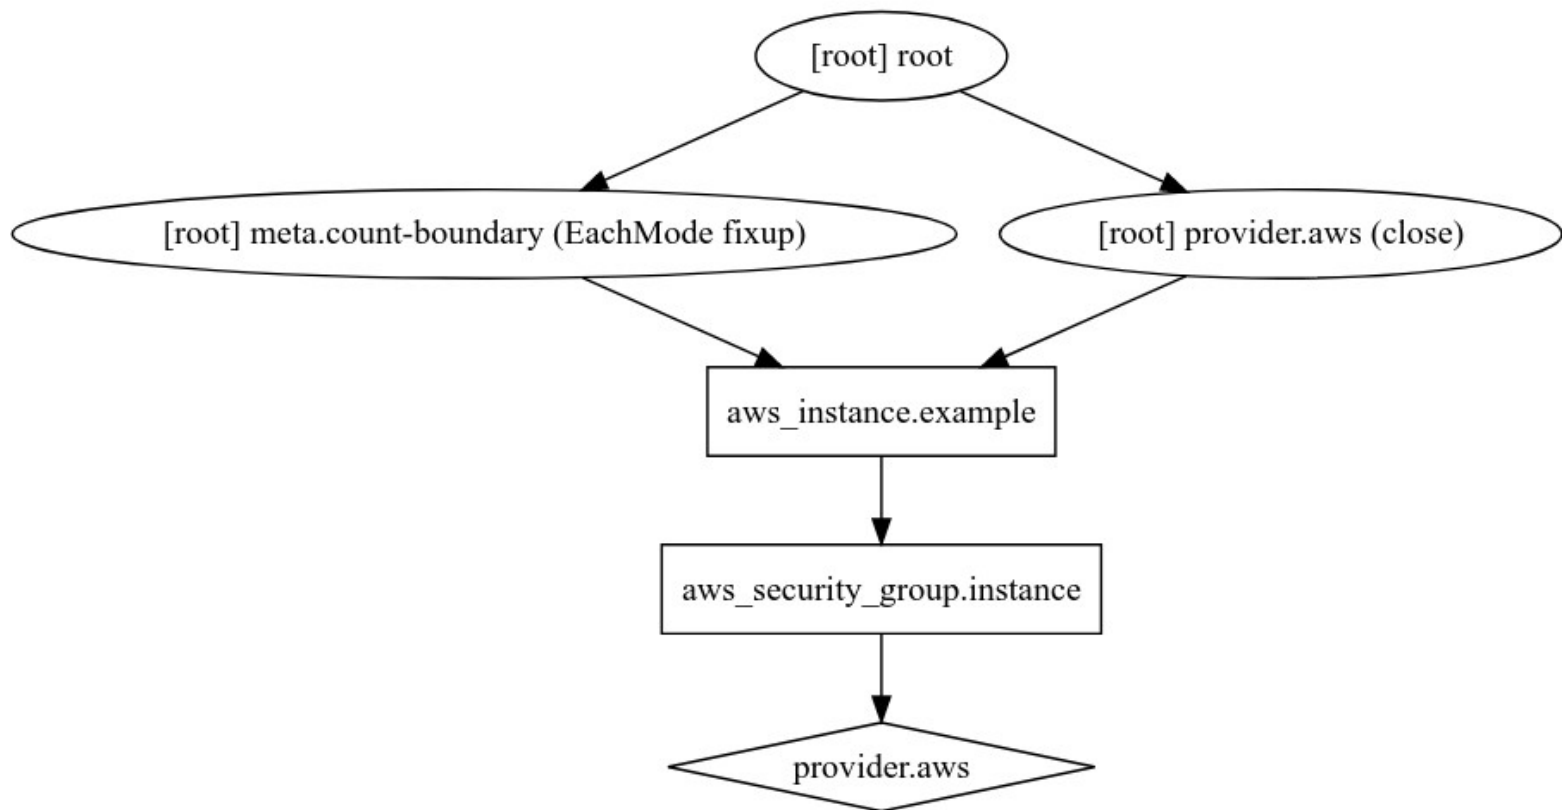
```
1 terraform graph
```

# Terraform Graph Output

```
digraph {
  compound = "true"
  newrank = "true"
  subgraph "root" {
    "[root] aws_instance.example" [label = "aws_instance.example", shape = "box"]
    "[root] aws_security_group.instance" [label = "aws_security_group.instance", shape = "box"]
    "[root] provider.aws" [label = "provider.aws", shape = "diamond"]
    "[root] aws_instance.example" -> "[root] aws_security_group.instance"
    "[root] aws_security_group.instance" -> "[root] provider.aws"
    "[root] meta.count-boundary (EachMode fixup)" -> "[root] aws_instance.example"
    "[root] provider.aws (close)" -> "[root] aws_instance.example"
    "[root] root" -> "[root] meta.count-boundary (EachMode fixup)"
    "[root] root" -> "[root] provider.aws (close)"
  }
}
```

# Terraform Graph Visual

- ◆ Use a desktop app such as Graphviz or



# NETWORK SECURITY

- ◆ All our example deploy not only into your Default VPC (as mentioned earlier), but also the default subnets of that VPC
- ◆ Running a server in a public subnet is fine for a quick experiment, but in real-world usage, it's a security risk
- ◆ For production systems, you should deploy all of your servers, and certainly all of your data stores, in private subnets



# Using Variables

**Using Variables**  
Adding Scalability

# Deploy a Configurable Web Server

- ◆ Don't Repeat Yourself (DRY) principle
- ◆ However, we violated it
  - the web server port 8080 is duplicated in both the security group and the User Data configuration
- ◆ So, DRY:
  - every piece of knowledge must have a single, unambiguous, authoritative representation within a system



# Terraform Input Variables

```
1 variable "NAME" {  
2   [CONFIG ...]  
3 }
```

- ◆ description
  - It's always a good idea to use this parameter to document how a variable is used
- ◆ default, or use these ways:
  - passing it in at the command line (using the -var option)
  - via a file (using the -var-file option)
  - via an environment variable
- ◆ type
  - enforce type constraints on the variables a user passes in
  - type constraints: string, number, bool, list, map, set, object, tuple, and any

# Examples of Terraform Input Variables

- ◆ Input variable that checks that the value you pass in is a number:

```
1 variable "number_example" {  
2     description = "An example of a number variable"  
3     type        = number  
4     default     = 42  
5 }
```



# Examples of Terraform Input Variables

- ◆ List input variable with all numbers

```
1 variable "list_numeric_example" {  
2     description = "An example of a numeric list"  
3     type        = list(number)  
4     default     = [1, 2, 3]  
5 }
```

# Examples of Terraform Input Variables

- ◆ A map of strings

```
1 variable "map_example" {  
2   description = "An example of a map  
3   type        = map(string)  
4  
5   default = {  
6     key1 = "value1"  
7     key2 = "value2"  
8     key3 = "value3"  
9   }  
10 }
```

# Quiz

- ◆ Consider the following Terraform 0.12 configuration snippet. How would you define the `cidr_block` for `us-east-1` in the `aws_vpc` resource using a variable?

```
1 variable "vpc_cidrs" {  
2   type = map  
3   default = {  
4     us-east-1 = "10.0.0.0/16"  
5     us-east-2 = "10.1.0.0/16"  
6     us-west-1 = "10.2.0.0/16"  
7     us-west-2 = "10.3.0.0/16"  
8   }  
9 }  
10  
11 resource "aws_vpc" "shared" {  
12   cidr_block = _____  
13 }
```

- ◆ A. `var.vpc_cidrs["us-east-1"]`
- ◆ B. `var.vpc_cidrs.0`
- ◆ C. `vpc_cidrs["us-east-1"]`
- ◆ D. `var.vpc_cidrs[0]`

# Quiz

- ◆ You have defined the values for your variables in the file `terraform.tfvars`, and saved it in the same directory as your Terraform configuration. Which of the following commands will use those values when creating an execution plan?
  - A. `terraform plan`
  - B. `terraform plan -var-file=terraform.tfvars`
  - C. All of the above
  - D. None of the above

# OO Coding with Terraform!

```
1 variable "object_example" {
2   description = "An example of a structural type"
3   type       = object({
4     name      = string
5     age       = number
6     tags      = list(string)
7     enabled   = bool
8   })
9
10  default = {
11    name      = "value1"
12    age       = 42
13    tags      = ["a", "b", "c"]
14    enabled   = true
15  }
16 }
```

# "server\_port" Variable

```
1 variable "server_port" {  
2   description = "Server port for HTTP requests"  
3   type        = number  
4 }
```

# Using "server\_port" Variable

- ◆ If you run `terraform apply`, you will get this message:

```
1 var.server_port
2   "Server port for HTTP requests"
3   Enter a value:
```

- ◆ Your choices now are:
  - Enter a value :)
  - `terraform plan -var "server_port=8080"`
  - `export TF_VAR_server_port=8080`
  - Supply a default

# How to Use Your Variable

- ◆ Simply, use `var` , like this: `var.<VARIABLE_NAME>`
- ◆ For example

```
1 resource "aws_security_group" "instance" {  
2     name = "terraform-example-instance"  
3  
4     ingress {  
5         from_port    = var.server_port  
6         to_port      = var.server_port  
7         protocol     = "tcp"  
8         cidr_blocks  = ["0.0.0.0/0"]  
9     }  
10 }
```



# New Expression - Interpolation

- ◆ "\${...}"
- ◆ Now, let us use the same `server_port` inside of `User Data`

```
1 user_data = <<-EOF
2   #!/bin/bash
3   echo "Hello, World" > index.html
4   nohup busybox httpd -f -p ${var.server_port} &
5   EOF
```

# Setting an "output" variable

- ◆ Additional variables
- ◆ description
  - It is always a good idea to document
- ◆ sensitive
  - true will instruct Terraform not to log this output at the end of terraform apply
  - For sensitive material or secrets such as passwords or private keys

```
1 output "<NAME>" {  
2   value = <VALUE>  
3   [CONFIG ...]  
4 }
```

# Output Variable For Our Script

```
1 output "public_ip" {  
2     value      = aws_instance.example.public_ip  
3     description = "The public IP address of the web server"  
4 }
```

# Adding Scalability

Using Variables  
**Adding Scalability**

# Motivation for a Cluster

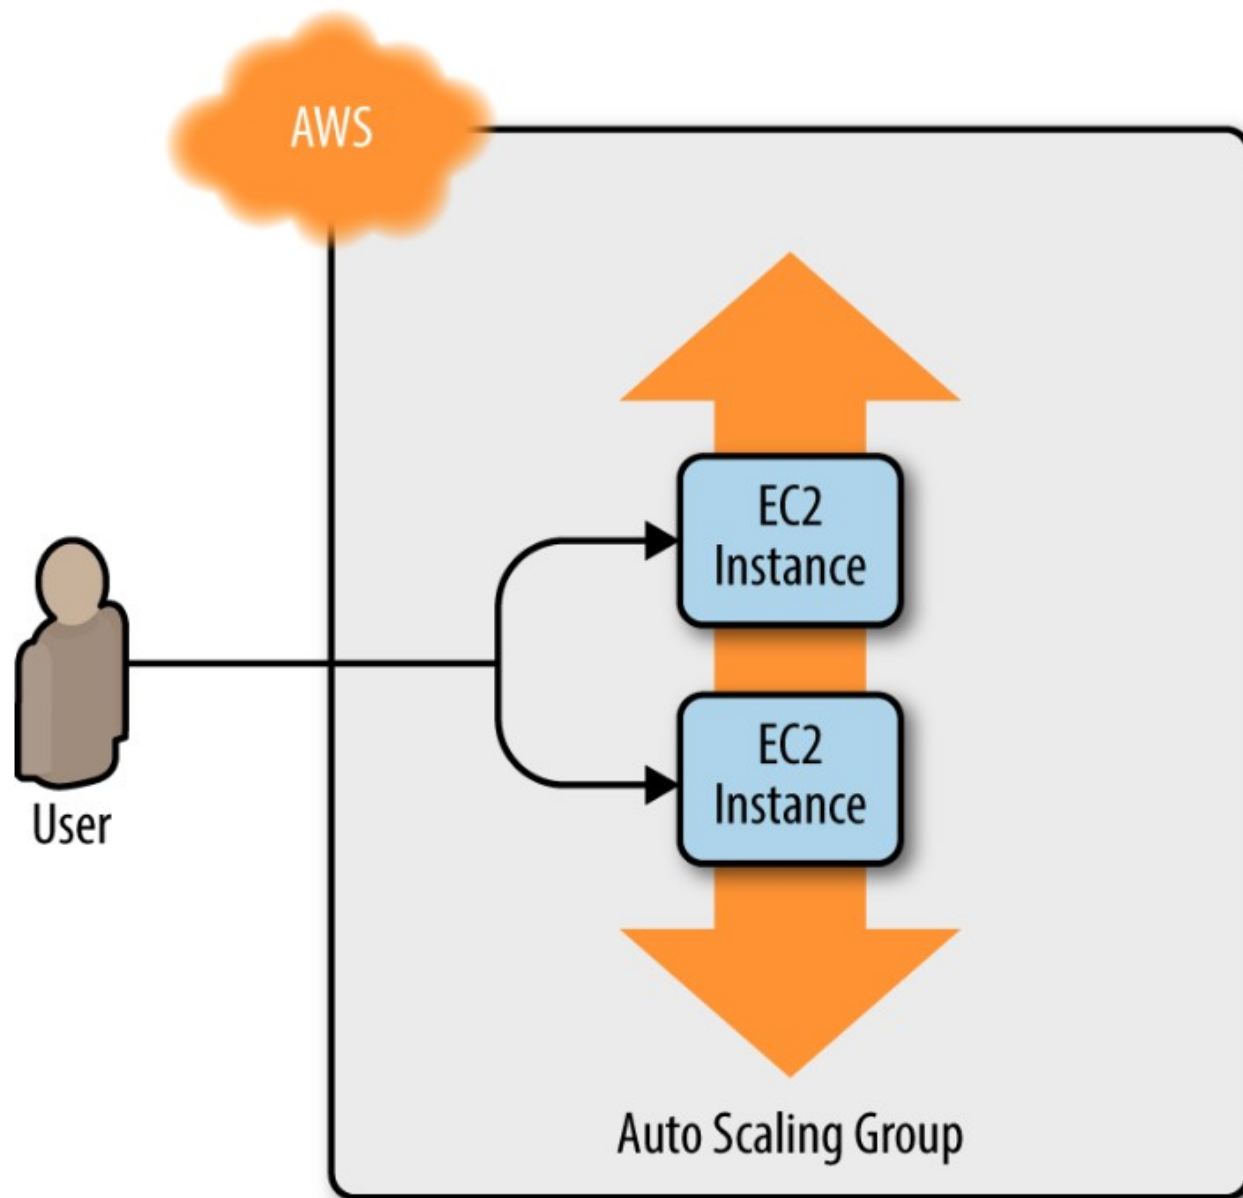
- ◆ Running a single server is a good start, but in the real world, a single server is a single point of failure
- ◆ If that server crashes, or if it becomes overloaded from too much traffic, users will be unable to access your site
- ◆ The solution is to run a cluster of servers, routing around servers that go down, and adjusting the size of the cluster up or down based on traffic
- ◆ We will need the following
  - Auto-scaling group (ASG)
  - VPC
  - Load balancer



# Discussion

- ◆ When clusters makes sense and when it is not worth it?
- ◆ Imagine a parallel research app for a cluster that would be rearchitected for Cloud
  - would have instances
  - a task dispatcher, etc.
  - would have to manage those instances, etc.
- ◆ Alternatives
  - Very large instance with 100+ CPUs
  - Memory pool from Kove

# Auto-Scaling Group (ASG)



# ASG Described in Terraform

- ◆ To create an ASG, first describe the instance that goes into it
  - Create a launch configuration
  - The `aws_launch_configuration` resource
    - uses almost exactly the same parameters as the `aws_instance` resource
    - `ami` is now `image_id`
    - `vpc_security_group_ids` is now `security_groups`
    - put this instead of

```
1 resource "aws_launch_configuration" "example" {
2     image_id      = "ami-0c55b159cbfafa1f0"
3     instance_type = "t2.micro"
4     security_groups = [aws_security_group.instance.id]
5
6     user_data = <<-EOF
7     #!/bin/bash
8         echo "Hello, World" > index.html
9         nohup busybox httpd -f -p ${var.server_port} &
10        EOF
11 }
```



# The ASG Itself with "aws\_autoscaling\_group"

- ◆ ASG will run between 2 and 10 instances
- ◆ each tagged with the name terraform-asg-example
- ◆ ASG uses a reference to fill in the launch configuration name

```
1 resource "aws_autoscaling_group" "example" {  
2     launch_configuration =  
3         aws_launch_configuration.example.name  
4  
5     min_size = 2  
6     max_size = 10  
7  
8     tag {  
9         key           = "Name"  
10        value        = "terraform-asg-example"  
11        propagate_at_launch = true  
12    }  
13 }
```

# "Lifecycle" Setting

- ◆ Use `create_before_destroy`
- ◆ If you set `create_before_destroy` to `true`
  - Terraform will invert the order in which it replaces resources
  - create the replacement resource first
  - then deleting the old resource

```
1 resource "aws_launch_configuration" "example" {
2     image_id      = "ami-0c55b159cbfaffe1f0"
3     instance_type = "t2.micro"
4     security_groups = [aws_security_group.instance.id]
5
6     user_data = <<-EOF
7         #!/bin/bash
8         echo "Hello, World" > index.html
9         nohup busybox httpd -f -p ${var.server_port} &
10        EOF
11     # Required when launching configuration with an ASG
12     lifecycle {
13         create_before_destroy = true
14     }
15 }
```

## One More Parameter: "subnet\_ids"

- ◆ specifies to the ASG into which VPC subnets the EC2 Instances should be deployed
- ◆ Each subnet lives in an isolated AWS AZ
- ◆ By deploying your instances across multiple subnets
  - you make it fault-tolerant
- ◆ Instead of hard-coding the list of subnet, we will get them from data sources

# Data Sources

- ◆ data source a piece of read-only information that is fetched from the provider (in this case, AWS) every time you run Terraform
- ◆ data source in your configurations is a way to query the provider's APIs for data
- ◆ AWS data sources include
  - VPC data
  - subnet data
  - AMI IDs
  - IP address ranges
  - more

# Data Source Syntax

```
1 data "<PROVIDER>_<TYPE>" "<NAME>" {  
2   [CONFIG ...]  
3 }
```

- ◆ Example: Do I have the default VPC?

```
1 data "aws_vpc" "default" {  
2   default = true  
3 }
```

# Getting data for "data source"

- ◆ Syntax
- ◆ `data.<PROVIDER>_<TYPE>.<NAME>.<ATTRIBUTE>`
- ◆ Example
- ◆ `data.aws_vpc.default.id`
- ◆ With this, you can find out the default subnet id

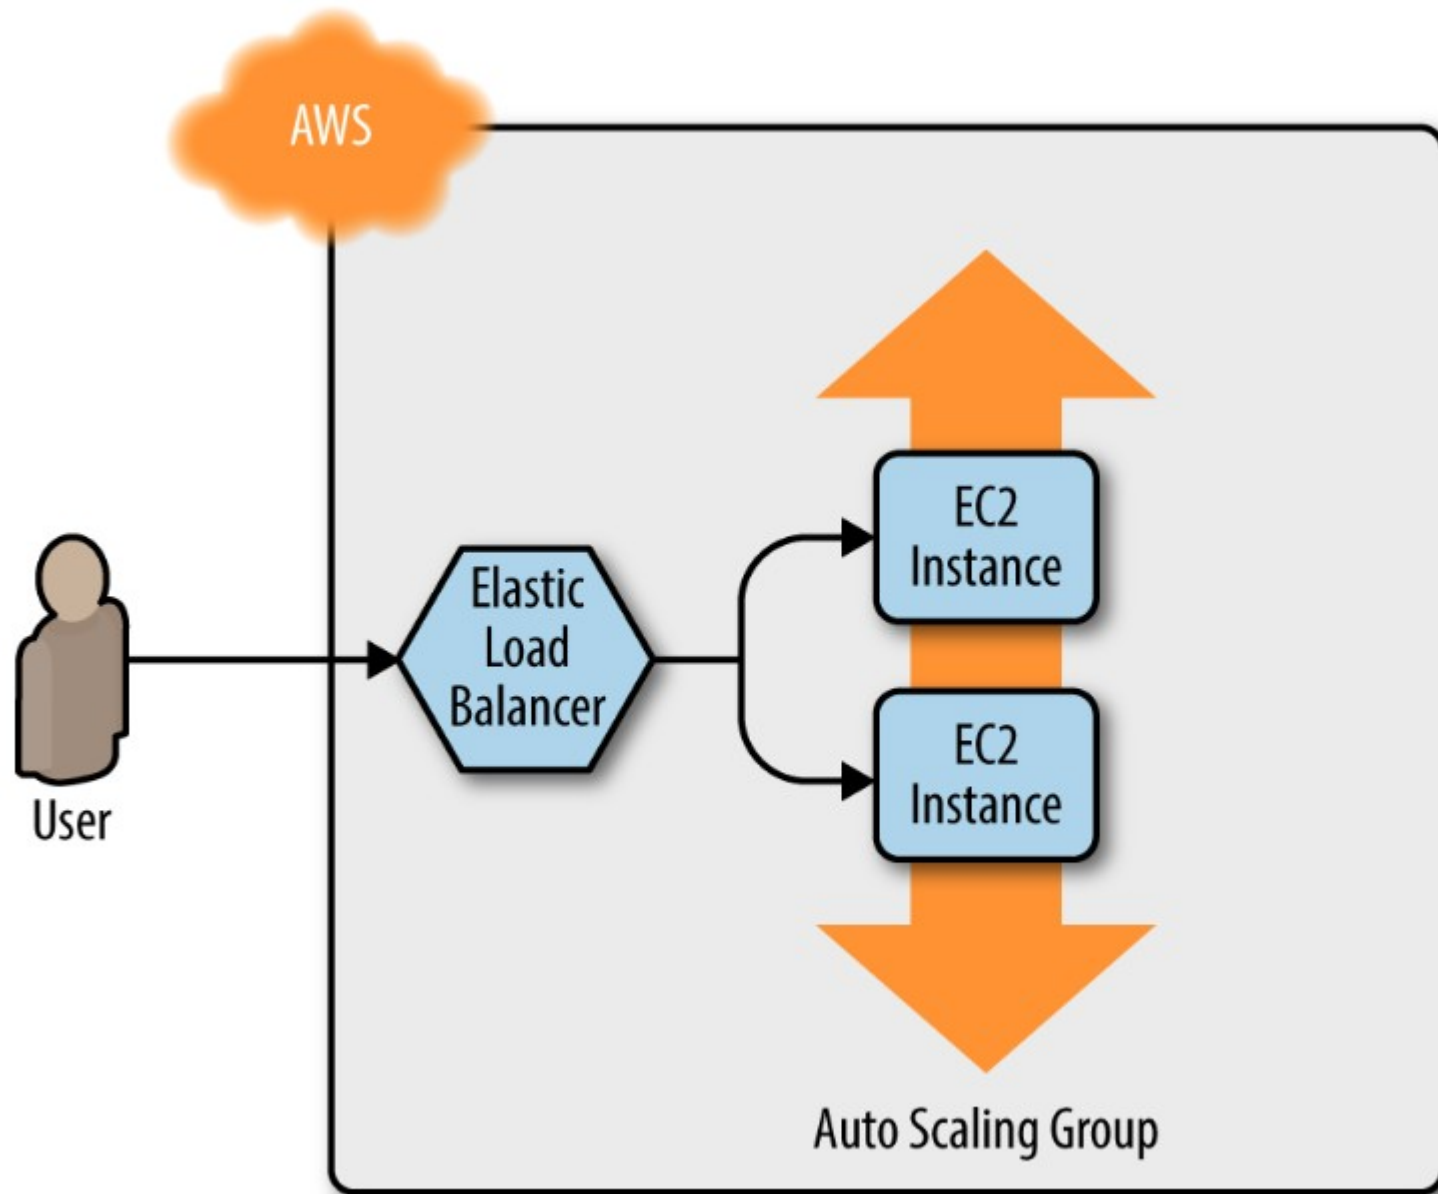
```
1 data "aws_subnet_ids" "default" {  
2   vpc_id = data.aws_vpc.default.id  
3 }
```

# Use the Default Subnet id

- ◆ Pull the subnet IDs out of the `aws_subnet_ids` data source
- ◆ Tell your ASG to use those subnets via the `vpc_zone_identifier` argument

```
1 resource "aws_autoscaling_group" "example" {
2   launch_configuration = aws_launch_configuration.example.name
3   vpc_zone_identifier  = data.aws_subnet_ids.default.ids
4
5   min_size = 2
6   max_size = 10
7
8   tag {
9     key           = "Name"
10    value          = "terraform-asg-example"
11    propagate_at_launch = true
12  }
13 }
```

# Load Balancer





# Using Load Balancer

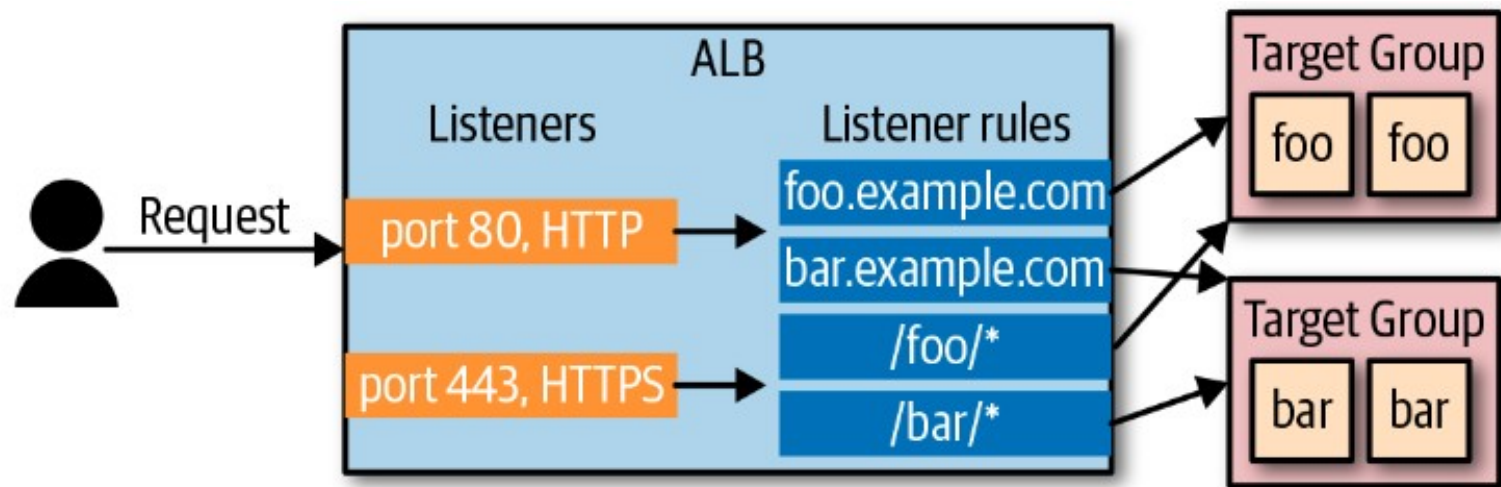
- ◆ Problem
  - you now have multiple servers, each with its own IP address
  - you want to give of your end users only a single IP to use
- ◆ Solution
  - deploy a load balancer to distribute traffic across your servers
- ◆ Advantage
  - highly available and scalable
- ◆ ELB to the rescue



# Load Balancer Types

- ◆ Application Load Balancer (ALB)
  - Best suited for load balancing of HTTP and HTTPS traffic
- ◆ Network Load Balancer (NLB)
  - Best suited for load balancing of TCP, UDP, and TLS traffic. Can scale up and down in response to load faster than the ALB (the NLB is designed to scale to tens of millions of requests per second). Operates at the transport layer (Layer 4) of the OSI model.
- ◆ Classic Load Balancer (CLB)
  - This is the “legacy” load balancer that predates both the ALB and NLB. It can handle HTTP, HTTPS, TCP, and TLS traffic, but with far fewer features than either the ALB or NLB. Operates at both the application layer (L7) and transport layer (L4) of the OSI model.

# Application Load Balancer (ALB)



# ALB Configuration

```
1 resource "aws_lb" "example" {  
2     name           = "terraform-asg-example"  
3     load_balancer_type = "application"  
4     subnets       = data.aws_subnet_ids.default.ids  
5 }
```

# ALB Listener

```
1 resource "aws_lb_listener" "http" {
2   load_balancer_arn = aws_lb.example.arn
3   port              = 80
4   protocol          = "HTTP"
5
6   # By default, return a simple 404 page
7   default_action {
8     type = "fixed-response"
9
10    fixed_response {
11      content_type = "text/plain"
12      message_body = "404: page not found"
13      status_code  = 404
14    }
15  }
16 }
```

# Security Group for ALB

```
1 resource "aws_security_group" "alb" {
2   name = "terraform-example-alb"
3
4   # Allow inbound HTTP requests
5   ingress {
6     from_port    = 80
7     to_port      = 80
8     protocol     = "tcp"
9     cidr_blocks  = ["0.0.0.0/0"]
10  }
11
12  # Allow all outbound requests
13  egress {
14    from_port    = 0
15    to_port      = 0
16    protocol     = "-1"
17    cidr_blocks  = ["0.0.0.0/0"]
18  }
19 }
```

# "aws\_lb resource" to Use Our Security Group

```
1 resource "aws_lb" "example" {  
2     name                = "terraform-asg-example"  
3     load_balancer_type  = "application"  
4     subnets            = data.aws_subnet_ids.default.ids  
5     security_groups     = [aws_security_group.alb.id]  
6 }
```

# Limits for Your ASG

```
1 resource "aws_lb_target_group" "asg" {
2     name      = "terraform-asg-example"
3     port      = var.server_port
4     protocol  = "HTTP"
5     vpc_id    = data.aws_vpc.default.id
6
7     health_check {
8         path      = "/"
9         protocol   = "HTTP"
10        matcher    = "200"
11        interval   = 15
12        timeout     = 3
13        healthy_threshold = 2
14        unhealthy_threshold = 2
15    }
16 }
```



# What the Target Group Do?

- ◆ health check your Instances by periodically sending an HTTP request to each Instance
- ◆ will consider the Instance “healthy” only if the Instance returns a response that matches the configured matcher
- ◆ we told the matcher to look for a 200 OK response the target group will automatically stop sending traffic to unhealthy instance

# Target Group Knows Its EC2 Instances

```
1 resource "aws_autoscaling_group" "example" {
2   launch_configuration = aws_launch_configuration.example.name
3   vpc_zone_identifier  = data.aws_subnet_ids.default.ids
4
5   target_group_arns = [aws_lb_target_group.asg.arn] # HERE
6   health_check_type = "ELB" # HERE
7
8   min_size = 2
9   max_size = 10
10
11   tag {
12     key           = "Name"
13     value         = "terraform-asg-example"
14     propagate_at_launch = true
15   }
16 }
```


# ALB Listener Rule

```
1 resource "aws_lb_listener_rule" "asg" {
2     listener_arn = aws_lb_listener.http.arn
3     priority     = 100
4
5     condition {
6         field = "path-pattern"
7         values = ["*"]
8     }
9
10    action {
11        type = "forward"
12        target_group_arn = aws_lb_target_group.asg.arn
13    }
14 }
```


# New Output - The DNS Name of the ALB


```
1 output "alb_dns_name" {  
2     value      = aws_lb.example.dns_name  
3     description = "The domain name of the load balancer"  
4 }
```






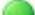


# Results of Upcoming Lab - Instances

Launch Instance 

Connect

Actions 

 Filter by tags and attributes or search by keyword

<input type="checkbox"/>	Name 	Instance ID 	Instance Type 	Availability Zone 	Instance State 
<input type="checkbox"/>	terraform-asg-example	i-060fd2fe038c6f2d0	t2.micro	us-east-2a	 running
<input type="checkbox"/>	terraform-example	i-0db5cc1490522774b	t2.micro	us-east-2c	 terminated
<input type="checkbox"/>	terraform-asg-example	i-0e34c1427a9e52c1f	t2.micro	us-east-2c	 running

# Results of Upcoming Lab - Load Balancer

Create Load Balancer

Actions ▾

Filter by tags and attributes or search by keyword					
<input type="checkbox"/>	Name ▴	DNS name ▾	State ▾	VPC ID ▾	Availability Zones ▾
<input type="checkbox"/>	terraform-asg-example	terraform-asg-example-1482...	active	vpc-f6d6bf9f	us-east-2a, us-east-2b,...

# Results of Upcoming Lab - Target Group

Create target group

Actions ▾

🔍 Filter by tags and attributes or search by keyword

<input type="checkbox"/>	Name ▴	Port ▾	Protocol ▾	Target type ▾	Load Balanc▾	VPC ID
<input type="checkbox"/>	terraform-asg-example	8080	HTTP	instance		vpc-f6d6bf9f