

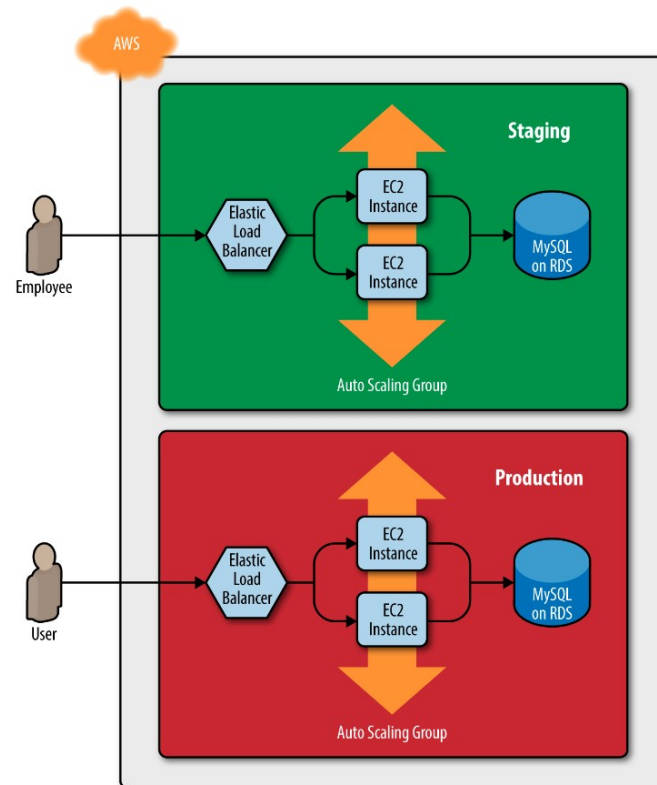
Reusable Infrastructure with Modules

The Plan

- ◆ Module basics
- ◆ Module inputs
- ◆ Module locals
- ◆ Module outputs
- ◆ Module gotchas
- ◆ Module versioning

Multiple Environments

- ◆ Cloud advantage: create multiple copies of the same environment
 - Production, Development, Test
 - Environments need to be similar if not identical
- ◆ We don't want to be able to re-use Terraform code across environments
 - DRY Principle: "Do not repeat yourself"

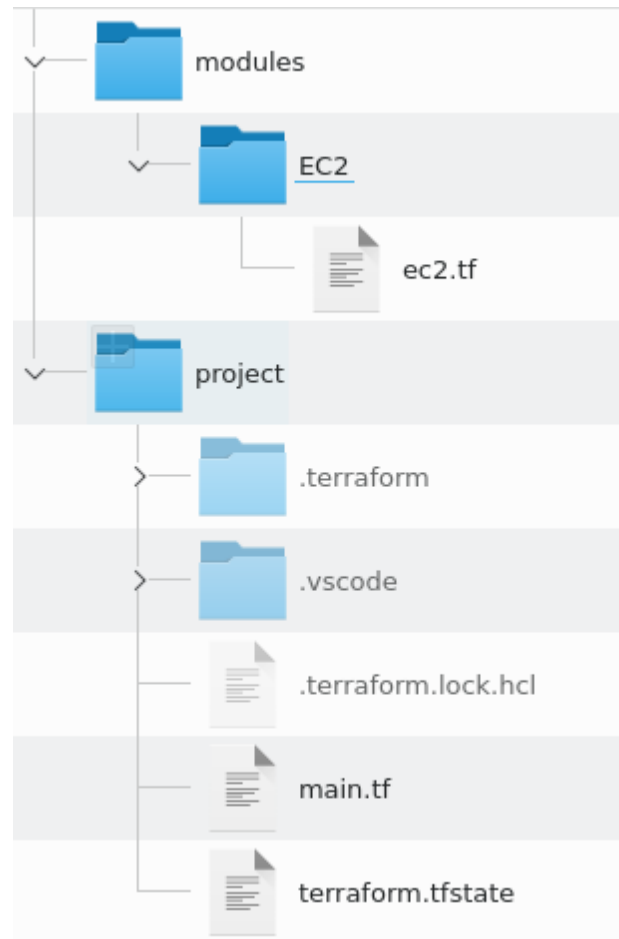


Module Basics

- ◆ Any folder containing Terraform files is a module
 - There are no special declarations or syntax required
 - Modules are containers for multiple resources that are used together
 - Modules are the primary strategy used to package and reuse Terraform resources
- ◆ Every Terraform configuration has at least one module
 - It is referred to as the "root" module
 - It consists of the Terraform files in the main working directory
- ◆ Modules (usually the root) may import other or "call" other modules
 - Modules that are being called are called "child" modules

Calling Modules Example

- ◆ If we are creating the same resource in multiple configurations, we can put it into a module
 - In this example, the demo project uses a module in the `modules/EC2` folder to create an EC2 instance
 - The folder structure looks like this:



Calling Modules Example

- ◆ The EC2 Module code is familiar

```
1 resource "aws_instance" "alpha" {  
2   ami = "ami-047a51fa27710816e"  
3   instance_type = "t2.micro"  
4   tags = {  
5     source = "EC2 Module"  
6   }  
}
```

- ◆ Calling it as a module is straightforward

```
1 module "EC2Defs" {  
2   source = "../modules/EC2"  
3 }
```

- ◆ The problem is that this module is not easily reusable because the `ami` and `instance_type` are hard-coded into the module code
 - We need to parameterize the module to make it reusable

Module Inputs

- ◆ Following the example of calling functions in a programming language, we want to be able to pass values to a module as parameters
- ◆ In the webserver example, we want to parameterize the code by adding three variables:

```
1 variable "cluster_name" {
2   description = "The name to use for all the cluster resources"
3   type        = string
4 }
5
6 variable "db_remote_state_bucket" {
7   description = "The name of the S3 bucket for the database's remote state"
8   type        = string
9 }
10
11 variable "db_remote_state_key" {
12   description = "The path for the database's remote state in S3"
13   type        = string
14 }
```

Module Parameters

- ◆ Now these variables can be used in the actual code instead of the hard-coded values

```
1 resource "aws_security_group" "alb" {
2   name = "${var.cluster_name}-alb"
3
4   ingress {
5     from_port    = 80
6     to_port      = 80
7     protocol     = "tcp"
8     cidr_blocks = ["0.0.0.0/0"]
9   }
10
11   egress {
12     from_port    = 0
13     to_port      = 0
14     protocol     = "-1"
15     cidr_blocks = ["0.0.0.0/0"]
16   }
17 }
```


Remote Back-end

- ◆ We can update the `terraform_remote_state` data source to use the `db_remote_state_bucket` and `db_remote_state_key` as its bucket and key parameter, respectively, to ensure we are reading the state file from the right environment:

```
1 data "terraform_remote_state" "db" {  
2   backend = "s3"  
3  
4   config = {  
5     bucket = var.db_remote_state_bucket  
6     key    = var.db_remote_state_key  
7     region = "us-east-2"  
8   }  
9 }
```

Staging Environment

- ◆ We can now pass arguments to these parameters in the call to the module `webserver_cluster`

```
1 module "webserver_cluster" {  
2   source = "../../modules/services/webserver-cluster"  
3  
4   cluster_name      = "webservers-stage"  
5   db_remote_state_bucket = "(YOUR_BUCKET_NAME)"  
6   db_remote_state_key   = "stage/data-stores/mysql/terraform.tfstate"  
7 }
```

- ◆ This module call can be made identically from
 - `stage/services/webserver-cluster/main` and
 - `prod/services/webserver-cluster/main.tf`

Module API

- ◆ This syntax for using input variables for a module by using the same syntax as setting arguments for a resource
- ◆ The input variables are the API of the module
 - Depending on the input variable arguments, each call to a module will result in a unique configuration
- ◆ Any place there is a hard-coded value in a module, we can replace it with a parameter
 - We want to choose the values to parameterize to produce maximum flexibility but balanced with ease of use

Parameterize the EC2 Cluster

- ◆ We start by defining the variable that will act as parameters

```
1 variable "instance_type" {  
2   description = "The type of EC2 Instances to run (e.g. t2.micro)"  
3   type        = string  
4 }  
5  
6 variable "min_size" {  
7   description = "The minimum number of EC2 Instances in the ASG"  
8   type        = number  
9 }  
10  
11 variable "max_size" {  
12   description = "The maximum number of EC2 Instances in the ASG"  
13   type        = number  
14 }
```

Parameterize the Module Code

- ◆ Replace the hard-coded values with variable references

```
1 resource "aws_launch_configuration" "example" {
2   image_id      = "ami-0c55b159cbfafelf0"
3   instance_type = var.instance_type
4   security_groups = [aws_security_group.instance.id]
5   user_data      = data.template_file.user_data.rendered
6
7   // Required when using a launch configuration with an auto scaling group.
8   // https://www.terraform.io/docs/providers/aws/r/launch_configuration.html
9
10  lifecycle {
11    create_before_destroy = true
12  }
13 }
```

Parameterize the ASG

- ◆ The ASG code can now be parameterized

```
1 resource "aws_autoscaling_group" "example" {
2   launch_configuration = aws_launch_configuration.example.name
3   vpc_zone_identifier  = data.aws_subnet_ids.default.ids
4   target_group_arns    = [aws_lb_target_group.asg.arn]
5   health_check_type    = "ELB"
6
7   min_size = var.min_size
8   max_size = var.max_size
9
10  tag {
11    key           = "Name"
12    value         = var.cluster_name
13    propagate_at_launch = true
14  }
15 }
```

Calling the Module - Revisited

- ◆ The `web_server` module can now be called with more parameters

```
1 module "webserver_cluster" {  
2   source = "../../modules/services/webserver-cluster"  
3  
4   cluster_name      = "webservers-stage"  
5   db_remote_state_bucket = "(YOUR_BUCKET_NAME)"  
6   db_remote_state_key   = "stage/data-stores/mysql/terraform.tfstate"  
7  
8   instance_type = "t2.micro"  
9   min_size      = 2  
10  max_size      = 2  
11 }
```

Module Reuse

- ◆ We can call the same module but with different parameters when we use it in a different situation

```
1 module "webserver_cluster" {
2   source = "../../modules/services/webserver-cluster"
3
4   cluster_name      = "webservers-prod"
5   db_remote_state_bucket = "(YOUR_BUCKET_NAME)"
6   db_remote_state_key   = "prod/data-stores/mysql/terraform.tfstate"
7
8   instance_type = "m4.large"
9   min_size      = 2
10  max_size      = 10
11 }
```


Module Locals

- ◆ There may exist hard-coded values that we want to convert to variables, to implement DRY for example
 - But when the module is called, the calling code should not be able to set the values of these variables
- ◆ Example, setting standard values that get repeated in code

```
1 resource "aws_lb_listener" "http" {
2   load_balancer_arn = aws_lb.example.arn
3   port              = 80
4   protocol          = "HTTP"
5
6   (...)
7 }
8
9 resource "aws_security_group" "alb" {
10  name = "${var.cluster_name}-alb"
11
12  ingress {
13    from_port = 80
14    to_port   = 80
15    (...)
16  }
17
18  (...)
19 }
```

Locals

- ◆ Local values can be defined in a `locals` block.
 - Work like variables
 - But they cannot be referenced or seen outside the module
 - Specifically, they are invisible to the calling code

```
1 locals {  
2   http_port      = 80  
3   any_port       = 0  
4   any_protocol   = "-1"  
5   tcp_protocol   = "tcp"  
6   all_ips        = ["0.0.0.0/0"]  
7 }
```

- ◆ And are referenced as `local`

```
1 resource "aws_lb_listener" "http" {  
2   load_balancer_arn = aws_lb.example.arn  
3   port              = local.http_port  
4   protocol          = "HTTP"  
5   (...)   
6 }
```

Module Outputs

- ◆ Modules can return values just like functions in programming languages
- ◆ This is done by defining output variables in the module
- ◆ For example, we can define the output variable `asg_name` in the `outputs.tf` file in the web-server cluster module

```
1 output "asg_name" {  
2     value      = aws_autoscaling_group.example.name  
3     description = "The name of the Auto Scaling Group"  
4 }
```

- ◆ We can then reference the value with the following syntax.
 - Reminder that the module name is the value we create when we *call* the module.

```
1 module.< MODULE_NAME >.< OUTPUT_NAME >  
2  
3 module.frontend.asg_name
```

Using Output Variables

- ◆ We can use the output variables from a module like any other variable
 - In the example, the name of the ASG is used to set an argument of another resource.

```
1 resource "aws_autoscaling_schedule" "scale_out_during_business_hours" {
2   scheduled_action_name = "scale-out-during-business-hours"
3   min_size              = 2
4   max_size              = 10
5   desired_capacity      = 10
6   recurrence             = "0 9 * * *"
7
8   autoscaling_group_name = module.webserver_cluster.asg_name
9 }
10
11 resource "aws_autoscaling_schedule" "scale_in_at_night" {
12   scheduled_action_name = "scale-in-at-night"
13   min_size              = 2
14   max_size              = 10
15   desired_capacity      = 2
16   recurrence             = "0 17 * * *"
17
18   autoscaling_group_name = module.webserver_cluster.asg_name
19 }
```

Variable Passthroughs

- ◆ One output variable can be "passed through" or used in a different output variable
- ◆ In the following example, the `dns_name` variable is defined in the `/modules/services/webserver-cluster/outputs.tf` file:

```
1 output "alb_dns_name" {  
2   value      = aws_lb.example.dns_name  
3   description = "The domain name of the load balancer"  
4 }
```

- ◆ This can be then "passed through" in the file `prod/services/webserver-cluster/outputs.tf`

```
1 output "alb_dns_name" {  
2   value      = module.webserver_cluster.alb_dns_name  
3   description = "The domain name of the load balancer"  
4 }
```

- ◆ Remember that all the dependencies between variables are resolved at planning time, not at apply time

Module Gotchas - Paths

- ◆ The hard-coded file paths are interpreted as relative to the current working directory
 - The problem is that this will not work if we are working with a module in a different directory
- ◆ To solve this issue, you can use an expression known as a path reference, which is of the form `path.<TYPE>`. Terraform supports the following types of path references:
 - `path.module` : Returns the file system path of the module where the expression is defined
 - `path.root` : Returns the file system path of the root module
 - `path.cwd` : Returns the file system path of the current working directory, usually the same as `path.root`

Module Path

- ◆ In this example, the template file is located with a path relative to the module, but if we hard-code the path, it will be interpreted as relative to the current working directory
 - By using the `path.module` construct, we ensure the file reference remains relative to the module

```
1 data "template_file" "user_data" {  
2   template = file("${path.module}/user-data.sh")  
3  
4   vars = {  
5     server_port = var.server_port  
6     db_address  = data.terraform_remote_state.db.outputs.address  
7     db_port     = data.terraform_remote_state.db.outputs.port  
8   }  
9 }
```

Module Gotcha - Inline Blocks

- ◆ The configuration for some Terraform resources can be defined either as inline blocks or as separate resources
 - When creating a module, you should always prefer using a separate resource
- ◆ Inline block example

```
1 resource "aws_security_group" "alb" {  
2   name = "${var.cluster_name}-alb"  
3  
4   ingress {  
5     from_port    = local.http_port  
6     to_port      = local.http_port  
7     protocol     = local.tcp_protocol  
8     cidr_blocks  = local.all_ips  
9   }  
10  
11  egress {  
12    from_port    = local.any_port  
13    to_port      = local.any_port  
14    protocol     = local.any_protocol  
15    cidr_blocks  = local.all_ips  
16  }  
17 }
```


Separate Resource

- ◆ You should change this module to define the exact same ingress and egress rules by using separate `aws_security_group_rule` resources

```
1 resource "aws_security_group" "alb" {
2   name = "${var.cluster_name}-alb"
3 }
4
5 resource "aws_security_group_rule" "allow_http_inbound" {
6   type           = "ingress"
7   security_group_id = aws_security_group.alb.id
8
9   from_port    = local.http_port
10  to_port      = local.http_port
11  protocol     = local.tcp_protocol
12  cidr_blocks  = local.all_ips
13 }
14
15 resource "aws_security_group_rule" "allow_all_outbound" {
16   type           = "egress"
17   security_group_id = aws_security_group.alb.id
18
19   from_port    = local.any_port
20   to_port      = local.any_port
21   protocol     = local.any_protocol
22   cidr_blocks  = local.all_ips
23 }
```

Inline Blocks

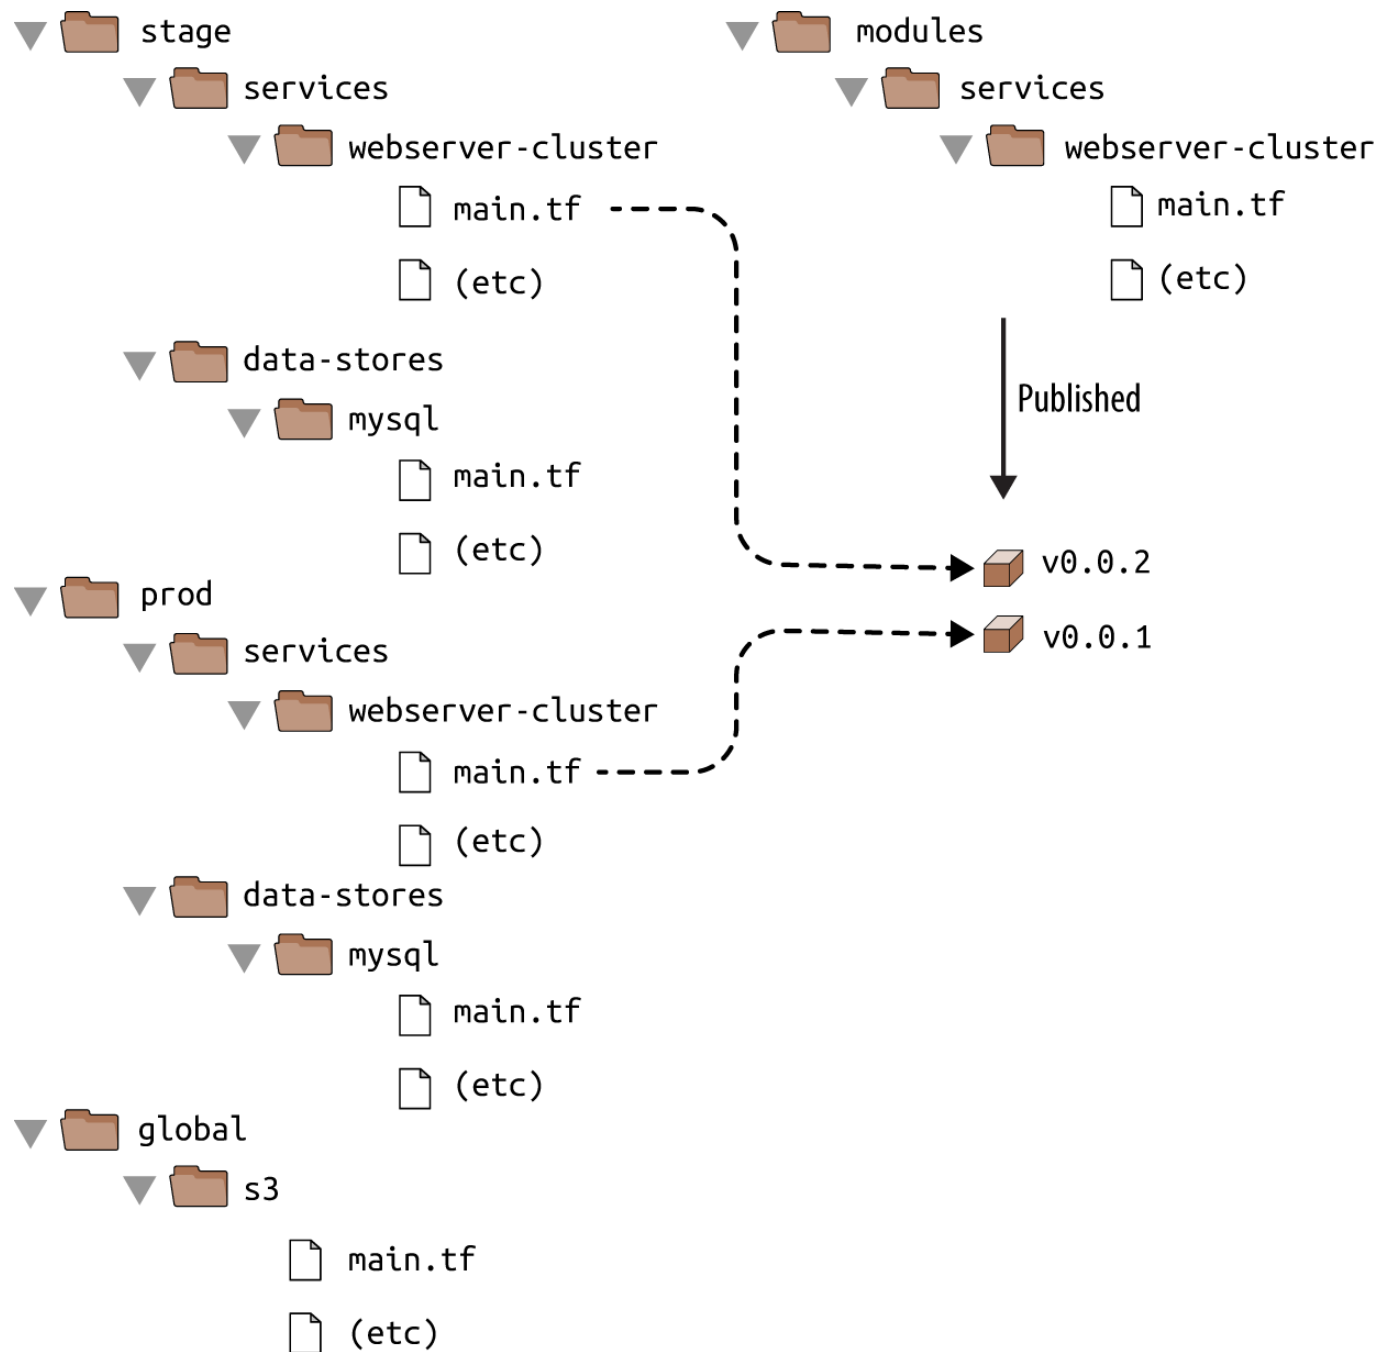
- ◆ Using a mix of inline blocks and separate resources may cause errors where routing rules conflict and overwrite one another
 - Use one or the other
 - When creating a module, you should always try to use a separate resource instead of the inline block
 - This allows for more flexible modules
- ◆ For example, changing a security group rule to allow a testing port is easier to do with a separate resource than having to edit inline blocks

```
1 resource "aws_security_group_rule" "allow_testing_inbound" {  
2     type                = "ingress"  
3     security_group_id = module.webserver_cluster.alb_security_group_id  
4  
5     from_port    = 12345  
6     to_port      = 12345  
7     protocol     = "tcp"  
8     cidr_blocks  = ["0.0.0.0/0"]  
9 }
```

Module Versioning

- ◆ If the staging and production environment point to the same module folder, any change in that folder will affect both environments on the very next deployment
 - This creates a coupling between environments and modules that can cause problems
- ◆ To solve this problem, we use a standard build management technique of using versions
 - As changes are made to a module, releases or versions of that module are published
 - Part of the configuration of any Terraform configuration plan is identification of which version of a module to include

Module Versioning Layout



Module Versioning

- ◆ An effective strategy is to use a repository tool like git and GitHub to publish releases of a module
 - Then the appropriate "release" of a module can be used

```
1 module "webserver_cluster" {
2   source = "github.com/foo/modules//webserver-cluster?ref=v0.0.1"
3
4   cluster_name      = "webservers-stage"
5   db_remote_state_bucket = "(YOUR_BUCKET_NAME)"
6   db_remote_state_key   = "stage/data-stores/mysql/terraform.tfstate"
7
8   instance_type = "t2.micro"
9   min_size      = 2
10  max_size      = 2
11 }
```

Semantic Versioning

- ◆ A common versioning scheme is "semantic versioning"
 - The format is MAJOR.MINOR.PATCH (e.g., 1.0.4)
 - There are specific rules on when you should increment each part of the version number
- ◆ MAJOR version increments when you make incompatible API changes
- ◆ MINOR version increments when you add functionality in a backward-compatible manner
- ◆ PATCH version increments when you make backward-compatible bug fixes