

Terraform Language Elements

Terraform Gotchas

Plan

- ◆ Loops
- ◆ If-Statements
- ◆ Deployment
- ◆ Gotchas
- ◆ Other language elements

希 福

Declarative Languages

- ◆ Declarative languages, like Terraform, normally do not have typical programming constructs like loops
- ◆ The challenge is expressing scenarios that require the conditional configuration of resources
 - For example, creating a module that creates resources only for certain users and not others
- ◆ Terraform primitives allow certain kinds of operations to allow dynamic and conditional configuration to be done
 - These do not look like standard constructs in programming languages that have the same functionality

Loops

- ◆ Terraform has several loop constructs to provide looping functionality in different scenarios
 - count parameter: to loop over resources
 - for_each expressions: to loop over resources and inline blocks within a functionality
 - for expressions: to loop over lists and maps
 - for string directive: to loop over lists and maps withing a string

Loops with

- ◆ The looping procedural code is implied and generated under the hood by Terraform
- ◆ We specify the number of iterations with the count, which usually would represent the number of copies of a resource
 - The following code creates three users with the names `neo.0`, `neo.1` and `neo.2`

```
1 resource "aws_iam_user" "example" {  
2     count = 3  
3     name  = "neo.${count.index}"  
4 }
```

- ◆ Under the hood something like this is conceptually happening:

```
1 resource "aws_iam_user" "example" {  
2     count = 3  
3     # for index = 0 to count  
4     name  = "neo.${count.index}"  
5 }
```

Array Look-ups

- ◆ We can supply list of value in arrays
 - Array elements can be reference with array notation
 - The length of the array (also strings and maps) returned from the built-in function `length()`

```
1 variable "user_names" {  
2   description = "Create IAM users with these names"  
3   type        = list(string)  
4   default     = ["neo", "trinity", "morpheus"]  
5 }
```

```
1 resource "aws_iam_user" "example" {  
2   count = length(var.user_names)  
3   name  = var.user_names[count.index]  
4 }
```

Arrays of Resources

- ◆ Using `count` on a resource creates an array of resources rather than just one resource
 - The lookup syntax is:
 - `<PROVIDER>_<TYPE>.<NAME>[INDEX].ATTRIBUTE`

```
1 output "neo_arn" {  
2   value      = aws_iam_user.example[0].arn  
3   description = "The ARN for user Neo"  
4 }
```

- ◆ To get all the users, a "splat" expression is used

```
1 output "all_arns" {  
2   value      = aws_iam_user.example[*].arn  
3   description = "The ARNs for all users"  
4 }
```

Arrays of Resources

- ◆ Running `terraform apply` will output the full array of resources

```
1 $ terraform apply
2
3 (...)
4
5 Apply complete! Resources: 3 added, 0 changed, 0 destroyed.
6
7 Outputs:
8
9 neo_arn = arn:aws:iam::123456789012:user/neo
10 all_arns = [
11     "arn:aws:iam::123456789012:user/neo",
12     "arn:aws:iam::123456789012:user/trinity",
13     "arn:aws:iam::123456789012:user/morpheus",
14 ]
```


Limitations of Count

- ◆ count can loop over resources but not inline blocks
- ◆ For example, we cannot iterate over the inline block for tag to generate multiple tag blocks dynamically

```
1 resource "aws_autoscaling_group" "example" {
2   launch_configuration = aws_launch_configuration.example.name
3   vpc_zone_identifier   = data.aws_subnet_ids.default.ids
4   target_group_arns     = [aws_lb_target_group.asg.arn]
5   health_check_type     = "ELB"
6
7   min_size = var.min_size
8   max_size = var.max_size
9
10  tag {
11    key           = "Name"
12    value         = var.cluster_name
13    propagate_at_launch = true
14  }
15 }
```

Limitations of Count

- ◆ Changing the values in a list modifies the created infrastructure
- ◆ If we create the infrastructure with this list:

```
1 variable "user_names" {  
2   description = "Create IAM users with these names"  
3   type        = list(string)  
4   default     = ["neo", "trinity", "morpheus"]  
5 }
```

- If `trinity` is removed then the correspondence between the array of resources and the list of names no longer is valid
- Terraform restores the mapping by recreating the resources

Limitations of Count

- ◆ Using

```
1 variable "user_names" {  
2   (...)  
3   default      = ["neo", "morpheus"]  
4 }
```

- ◆ Then terraform plan produces the following output

```
1 Terraform will perform the following actions:  
2  
3 # aws_iam_user.example[1] will be updated in-place  
4 ~ resource "aws_iam_user" "example" {  
5     id          = "trinity"  
6     ~ name      = "trinity" -> "morpheus"  
7 }  
8  
9 # aws_iam_user.example[2] will be destroyed  
10 - resource "aws_iam_user" "example" {  
11     - id          = "morpheus" -> null  
12     - name        = "morpheus" -> null  
13 }  
14  
15 Plan: 0 to add, 1 to change, 1 to destroy.
```

Loops with

- ◆ The `for_each` expression allows looping over lists, sets, and maps to create either:
 - multiple copies of an entire resource, or
 - multiple copies of an inline block within a resource
- ◆ The syntax is:

```
1 resource "< PROVIDER > < TYPE >" "< NAME >" {  
2   for_each = < COLLECTION >  
3  
4   [CONFIG ...]  
5 }
```

- ◆ The previous example is now:

```
1 variable "user_names" {  
2   description = "Create IAM users with these names"  
3   type        = list(string)  
4   default     = ["neo", "trinity", "morpheus"]  
5 }  
6  
7 resource "aws_iam_user" "example" {  
8   for_each = toset(var.user_names)  
9   name     = each.value  
10 }
```

Use of

- ◆ The function `toSet()` converts the `var.user_names` list into a set
 - `for_each` supports sets and maps only when used on a resource
- ◆ When `for_each` loops each user name is made available in the `each` value
 - The user name will also be available in `each.key`, but this is usually used only with maps of key/value pair.
- ◆ Once `for_each` is used on a resource, it creates a map of resources rather than array of resources
 - This is why we can't use a list with possible duplicates
 - This would lead to duplicate keys

Map Advantages

- ◆ Maps do not rely on position like lists do
 - Allows us to remove items from the middle of a collection safely
- ◆ Going back to the problem of deleting "trinity" with a map of resources we get:

```
1 terraform plan
2
3 Terraform will perform the following actions:
4
5 # aws_iam_user.example["trinity"] will be destroyed
6 - resource "aws_iam_user" "example" {
7     - arn                = "arn:aws:iam::123456789012:user/trinity" -> null
8     - name               = "trinity" -> null
9   }
10
11 Plan: 0 to add, 0 to change, 1 to destroy.
```

Inline Blocks with

◆ Setting some custom tags

```
1 variable "custom_tags" {
2   description = "Custom tags to set on the Instances in the ASG"
3   type        = map(string)
4   default     = {}
5 }
```

```
1 module "webserver_cluster" {
2   source = "../../../../../modules/services/webserver-cluster"
3   ...
4   cluster_name      = "webservers-prod"
5   db_remote_state_bucket = "(YOUR_BUCKET_NAME)"
6   db_remote_state_key   = "prod/data-stores/mysql/terraform.tfstate"
7   instance_type       = "m4.large"
8   min_size            = 2
9   max_size            = 10
10  ...
11  custom_tags = {
12    Owner      = "team-foo"
13    DeployedBy = "terraform"
14  }
15 }
```

Dynamic Inline Blocks

- ◆ To use a `for_each` to create a dynamic inline block, we use the following syntax and give an example of its use

```
1 dynamic "< VAR_NAME >" {  
2   for_each = < COLLECTION >  
3  
4   content {  
5     [CONFIG...]  
6   }  
7 }
```

```
1 resource "aws_autoscaling_group" "example" {  
2   launch_configuration = aws_launch_configuration.example.name  
3   ...  
4  
5   tag {  
6     key           = "Name"  
7     value         = var.cluster_name  
8     propagate_at_launch = true  
9   }  
10  
11   dynamic "tag" {  
12     for_each = var.custom_tags  
13  
14     content {  
15       key           = tag.key  
16       value         = tag.value  
17       propagate_at_launch = true  
18     }  
19   }  
20 }
```


Looping with Expressions

- ◆ Terraform allows operations on the data similar to operations in a programming language
- ◆ Syntax is:

```
1 [for < ITEM > in < LIST > : < OUTPUT >]
```

- ◆ Demonstrated in this code

```
1 variable "names" {  
2   description = "A list of names"  
3   type        = list(string)  
4   default     = ["neo", "trinity", "morpheus"]  
5 }  
6  
7 output "upper_names" {  
8   value = [for name in var.names : upper(name)]  
9 }  
10  
11 output "short_upper_names" {  
12   value = [for name in var.names : upper(name) if length(name) < 5]  
13 }
```

Working with Map Inputs

- ◆ The for expression can loop over a map as well

```
1 [for < KEY >, < VALUE > in < MAP > : < OUTPUT >]
```

- ◆ Example of use:

```
1 variable "hero_thousand_faces" {
2   description = "map"
3   type        = map(string)
4   default    = {
5     neo      = "hero"
6     trinity  = "love interest"
7     morpheus = "mentor"
8   }
9 }
10
11 output "bios" {
12   value = [for name, role in var.hero_thousand_faces : "${name} is the ${role}"]
13 }
```

```
1 // output produced by terraform apply
2 bios = [
3   "morpheus is the mentor",
4   "neo is the hero",
5   "trinity is the love interest",
6 ]
```

Outputting a Map

- ◆ Looping over a list or map can output a map using the syntax:

```
1 // Loop over a list and output a map
2 {for < ITEM > in < LIST > : < OUTPUT_KEY > => < OUTPUT_VALUE >}
```

- ◆

```
1 // Loop over a map and output a map
2 {for < KEY >, < VALUE > in < MAP > : < OUTPUT_KEY > => < OUTPUT_VALUE >}
```

- ◆ Revisiting the example

```
1 variable "hero_thousand_faces" {
2   description = "map"
3   type        = map(string)
4   default     = {
5     neo       = "hero"
6     trinity   = "love interest"
7     morpheus  = "mentor"
8   }
9 }
```

```
1 output "upper_roles" {
2   value = {for name, role in var.hero_thousand_faces : upper(name) => upper(role)}
3 }
```

```
1 upper_roles = {
2   "MORPHEUS" = "MENTOR"
3   "NEO"      = "HERO"
4   "TRINITY"  = "LOVE INTEREST"
5 }
```

Loops with the

- ◆ String directives allow for-loops and if-statements in strings using a syntax similar to string interpolations but instead of a dollar sign and curly braces (`${...}`), it uses a percent sign and curly braces (`%{...}`)
 - Terraform supports two types of string directives: for-loops and conditionals
- ◆ For loop syntax (collection is a list or map)

```
1 %{ for < ITEM > in < COLLECTION > }< BODY >%{ endfor }
```

String for loop Example

```
1 variable "names" {
2   description = "Names to render"
3   type        = list(string)
4   default     = ["neo", "trinity", "morpheus"]
5 }
```

```
1 output "for_directive" {
2
3   value = <<EOF
4   %{ for name in var.names }
5     ${name}
6   %{ endfor }
7   EOF
8 }
```

```
1 terraform apply
2
3 Apply complete! Resources: 0 added, 0 changed, 0 destroyed.
4
5 Outputs:
6
7 for_directive =
8   neo
9
10  trinity
11
12  morpheus
```

Trimming Whitespace

- ◆ A strip marker (~) in your string directive consumes all of the white space (spaces and newlines) either before the string directive (if the marker appears at the beginning of the string

```
1 output "for_directive_strip_marker" {  
2   value = <<EOF  
3   %{\~ for name in var.names }  
4   ${name}  
5   %{\~ endfor }  
6   EOF  
7 }
```

- ◆ Which produces:

```
1 for_directive_strip_marker =  
2 neo  
3 trinity  
4 morpheus
```

Conditionals

- ◆ There are also several different ways to do conditionals, each intended to be used in a slightly different scenario:
 - *count parameter* : Used for conditional resources
 - *for_each and for expressions* : Used for conditional resources and inline blocks within a resource
 - *if string directive* : Used for conditionals within a string

Conditionals with Count

- ◆ We can define a Boolean variable as our test condition:

```
1 variable "enable_autoscaling" {  
2   description = "If set to true, enable auto scaling"  
3   type        = bool  
4 }
```

- ◆ We can set the count on a resource to "0" which means that resource is not created

- ◆

```
1 < CONDITION > ? < TRUE_VAL > : < FALSE_VAL >
```

- ◆ Terraform allows ternary conditionals of the form:

- ◆ This allows for conditional creation of resources:

```
1 resource "aws_autoscaling_schedule" "scale_out_during_business_hours" {  
2   count = var.enable_autoscaling ? 1 : 0  
3   ...  
4 }
```


Conditional Count

- ◆ This Boolean can then be set when creating a `webserver_cluster`

```
1 module "webserver_cluster" {  
2   source = "../../../../../modules/services/webserver-cluster"  
3  
4   cluster_name      = "webservers-stage"  
5   db_remote_state_bucket = "(YOUR_BUCKET_NAME)"  
6   db_remote_state_key   = "stage/data-stores/mysql/terraform.tfstate"  
7  
8   instance_type      = "t2.micro"  
9   min_size            = 2  
10  max_size             = 2  
11  enable_autoscaling   = false  
12 }
```

Working with Non-Boolean

- ◆ The previous example worked because we could define a Boolean variable
 - However, we may have to decode information in a string to make a decision
- ◆ Example: We want to set a cloud-watch alarm that triggers when CPU credits are low
 - However, CPUcredits only "txxx" instances
 - Larger instance like m4.large do not return a CPU credit metric and will always appear to be in an INSUFFICIENT_DATA state
 - We want the metric to apply to only txxx instance but we don't want to create a special Boolean

Working with Non-boolean

- ◆ The solution is to utilize the fact that first letter of the instance type should be a "t"

```
1 resource "aws_cloudwatch_metric_alarm" "low_cpu_credit_balance" {  
2   count = format("%.1s", var.instance_type) == "t" ? 1 : 0  
3  
4   alarm_name = "${var.cluster_name}-low-cpu-credit-balance"  
5   namespace  = "AWS/EC2"  
6   metric_name = "CPUCreditBalance"  
7   ...  
8 }
```

- ◆ The format function to extract just the first character from `var.instance_type`.
 - If that character is a "t" (e.g., t2.micro), it sets the count to 1;
 - otherwise, it sets the count to 0
 - This way, the alarm is created only for instance types that actually have a CPUCreditBalance metric.

If-else Conditions

- ◆ There is no specific if-else construct but it can be emulated
- ◆ In the following example, webserver-cluster module pulls in the user-data.sh script via a template_file data source

```
1 data "template_file" "user_data" {  
2   template = file("${path.module}/user-data.sh")  
3  
4   vars = {  
5     server_port = var.server_port  
6     db_address  = data.terraform_remote_state.db.outputs.address  
7     db_port     = data.terraform_remote_state.db.outputs.port  
8   }  
9 }
```

```
1 #!/bin/bash  
2  
3 cat > index.html <<EOF  
4 <h1>Hello, World</h1>  
5 <p>DB address: ${db_address}</p>  
6 <p>DB port: ${db_port}</p>  
7 EOF
```

If-else Conditions

- ◆ If we want to roll out a different version as well:

```
1 data "template_file" "user_data_new" {
2   template = file("${path.module}/user-data-new.sh")
3
4   vars = {
5     server_port = var.server_port
6   }
7 }
```

- ◆ We can define a Boolean:

```
1 #!/bin/bash
2
3 echo "Hello, World, v2" > index.html
4 nohup busybox httpd -f -p ${server_port} &
```

```
1 variable "enable_new_user_data" {
2   description = "If set to true, use the new User Data script"
3   type        = bool
4 }
```

Emulation if-else

- ◆ We can use the count parameter and a conditional expression to emulate an if-else condition

```
1 resource "aws_iam_user_policy_attachment" "neo_cloudwatch_full_access" {  
2   count = var.give_neo_cloudwatch_full_access ? 1 : 0  
3  
4   user      = aws_iam_user.example[0].name  
5   policy_arn = aws_iam_policy.cloudwatch_full_access.arn  
6 }  
7  
8 resource "aws_iam_user_policy_attachment" "neo_cloudwatch_read_only" {  
9   count = var.give_neo_cloudwatch_full_access ? 0 : 1  
10  
11   user      = aws_iam_user.example[0].name  
12   policy_arn = aws_iam_policy.cloudwatch_read_only.arn  
13 }
```

Accessing Output

- ◆ The problem: we conditionally create a resource but need to access some output attribute on the resource
 - Example: what if you wanted to offer two different User Data scripts in the webserver-cluster module and allow users to pick which one is executed?
 - Currently, the webserver-cluster module pulls in the user-data.sh script via a template_file data source:

```
1 data "template_file" "user_data" {
2   template = file("${path.module}/user-data.sh")
3
4   vars = {
5     server_port = var.server_port
6     db_address  = data.terraform_remote_state.db.outputs.address
7     db_port     = data.terraform_remote_state.db.outputs.port
8   }
9 }
```

Accessing Output

- ◆ The current `user-data.sh` script looks like this:

```
1 #!/bin/bash
2 cat > index.html << EOF
3 <h1>Hello, World</h1>
4 <p>DB address: ${db_address}</p>
5 <p>DB port: ${db_port}</p>
6 EOF
7 nohup busybox httpd -f -p ${server_port} &
```

- ◆ We want to allow some server clusters to use this alternative script called `user-data-new.sh` :

```
1 #!/bin/bash
2 echo "Hello, World, v2" > index.html
3 nohup busybox httpd -f -p ${server_port} &
```

- ◆ To use this script, you need a new `template_file` data source:

```
1 data "template_file" "user_data_new" {
2   template = file("${path.module}/user-data-new.sh")
3
4   vars = {
5     server_port = var.server_port
6   }
7 }
```


Making Choices

- ◆ Add a new Boolean input variable in `modules/services/webserver-cluster/variables.tf`:

```
1 variable "enable_new_user_data" {  
2     description = "If set to true, use the new User Data script"  
3     type        = bool  
4 }
```

- ◆ Use the if-else-statement trick to ensure that only one of the `template_file` data sources is created:

```
1 data "template_file" "user_data" {  
2     count = var.enable_new_user_data ? 0 : 1  
3  
4     template = file("${path.module}/user-data.sh")  
5     ...  
6 }  
7  
8 data "template_file" "user_data_new" {  
9     count = var.enable_new_user_data ? 1 : 0  
10  
11     template = file("${path.module}/user-data-new.sh")  
12     ...  
13 }
```

Making Choices

- ◆ Set the `user_data` parameter of the `aws_launch_configuration` resource to the `template_file` that actually exists
 - The conditional checks to see if the old file actually exists by checking the length of the old file array - if the file exists, it will have a non-zero length
 - Because of lazy evaluation, the 0th index of only the actually existing list will be evaluated

```
1 resource "aws_launch_configuration" "example" {
2   image_id      = "ami-0c55b159cbfafa1f0"
3   instance_type = var.instance_type
4   security_groups = [aws_security_group.instance.id]
5
6   user_data = (
7     length(data.template_file.user_data[*]) > 0
8     ? data.template_file.user_data[0].rendered
9     : data.template_file.user_data_new[0].rendered
10  )
11  ...
12 }
```

Conditional Use Case

- ◆ The new User Data script can be used in the staging environment by setting the `enable_new_user_data` parameter to true in `live/stage/services/webserver-cluster/main.tf`:

```
1 module "webserver_cluster" {
2   source = "../../../modules/services/webserver-cluster"
3
4   cluster_name      = "webservers-stage"
5   db_remote_state_bucket = "(YOUR_BUCKET_NAME)"
6   db_remote_state_key   = "stage/data-stores/mysql/terraform.tfstate"
7
8   instance_type      = "t2.micro"
9   min_size           = 2
10  max_size            = 2
11  enable_autoscaling  = false
12  enable_new_user_data = true
13 }
```

Conditional Use Case

- ◆ The old version of the script by setting `enable_new_user_data` to `false` in `live/prod/services/webserver-cluster/main.tf`:

```
1 module "webserver_cluster" {
2   source = "../../../modules/services/webserver-cluster"
3
4   cluster_name      = "webservers-prod"
5   db_remote_state_bucket = "(YOUR_BUCKET_NAME)"
6   db_remote_state_key   = "prod/data-stores/mysql/terraform.tfstate"
7
8   instance_type      = "m4.large"
9   min_size           = 2
10  max_size            = 10
11  enable_autoscaling  = true
12  enable_new_user_data = false
13  custom_tags = {
14    Owner      = "team-foo"
15    DeployedBy = "terraform"
16  }
17 }
```

Conditionals with `for_each` and `for`

- ◆ If you pass a `for_each` expression an empty collection, it will produce 0 resources or 0 inline blocks
- ◆ if you pass it a nonempty collection, it will create one or more resources or inline blocks
- ◆ The only question is how can you conditionally decide if the collection should be empty or not?
 - The answer is to combine the `for_each` expression with the `for` expression
- ◆ Consider the example if setting dynamic tags

```
1 dynamic "tag" {  
2   for_each = var.custom_tags  
3  
4   content {  
5     key           = tag.key  
6     value         = tag.value  
7     propagate_at_launch = true  
8   }  
9 }
```

Conditionals with for_each and for

- ◆ The nested for expression loops over `var.custom_tags`
 - Converts each value to uppercase
 - Then uses a conditional in the for expression to filter out any key set to Name because the module already sets its own Name tag

```
1 dynamic "tag" {
2     for_each = {
3         for key, value in var.custom_tags:
4             key => upper(value)
5             if key != "Name"
6         }
7
8     content {
9         key                = tag.key
10        value               = tag.value
11        propagate_at_launch = true
12    }
13 }
```

Conditionals with the if String Directive

- ◆ The string conditional directive has the form

```
1 {% if < CONDITION > }< TRUEVAL >{% endif }
```

- ◆ For example

```
1 variable "name" {  
2   description = "A name to render"  
3   type       = string  
4 }  
5  
6 output "if_else_directive" {  
7   value = "Hello, {% if var.name != "" }${var.name}{% else }(unnamed){% endif }"  
8 }
```

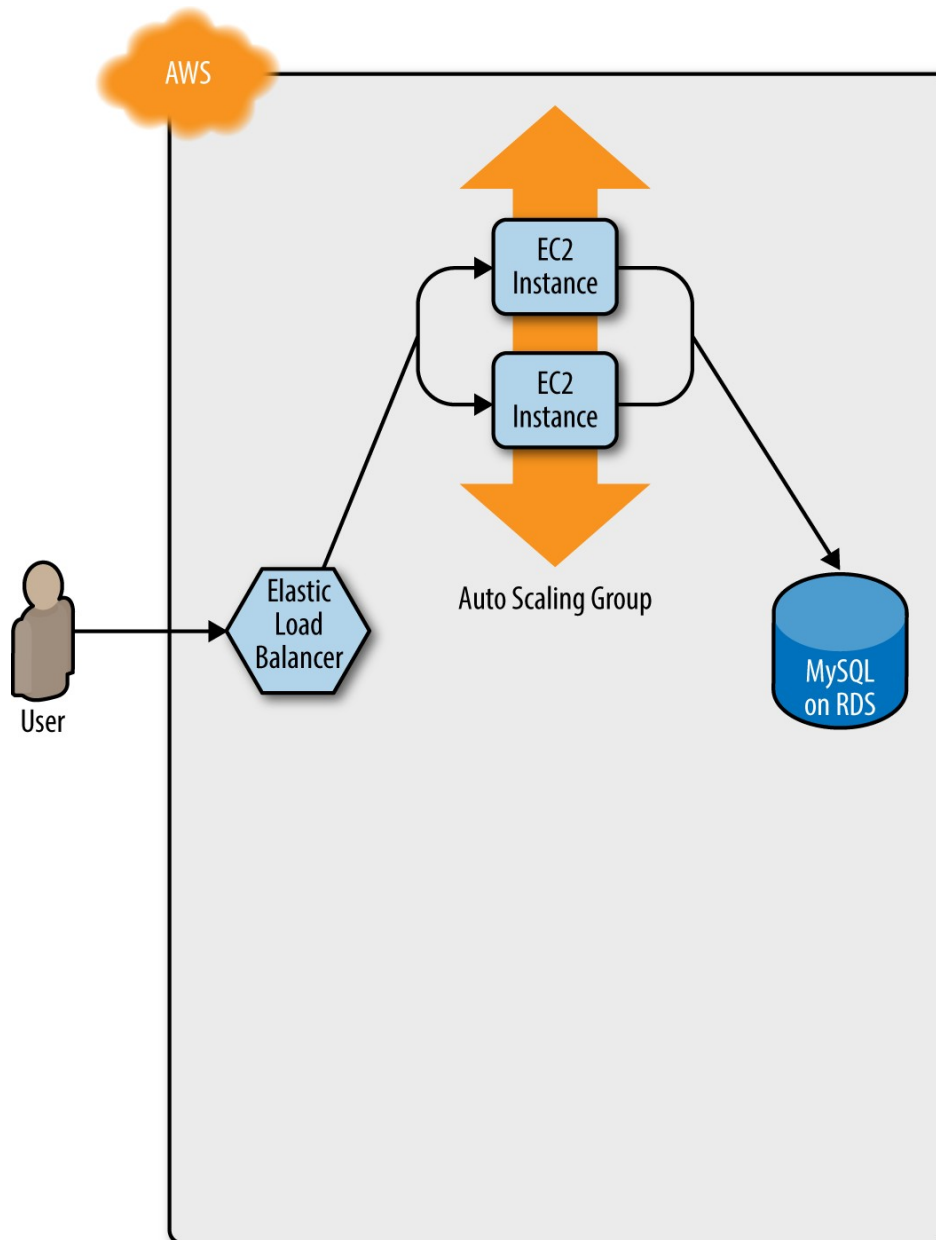
Zero-Downtime Deployment

- ◆ The challenge is to update a cluster without causing downtime for users
 - How do you deploy a new Amazon Machine Image (AMI) across the cluster?
- ◆ If we are deploying a new version of our app, we don't want there to be downtime as we switch over
- ◆ We full test and deploy our app in a test are to ensure it is working before we make the transition
- ◆ We then deploy the application into a new launch configuration which will be the target of the auto-scaling group
- ◆ However the challenge is switching launch configurations, if destroy the old one, then we have downtime while the new one is being created

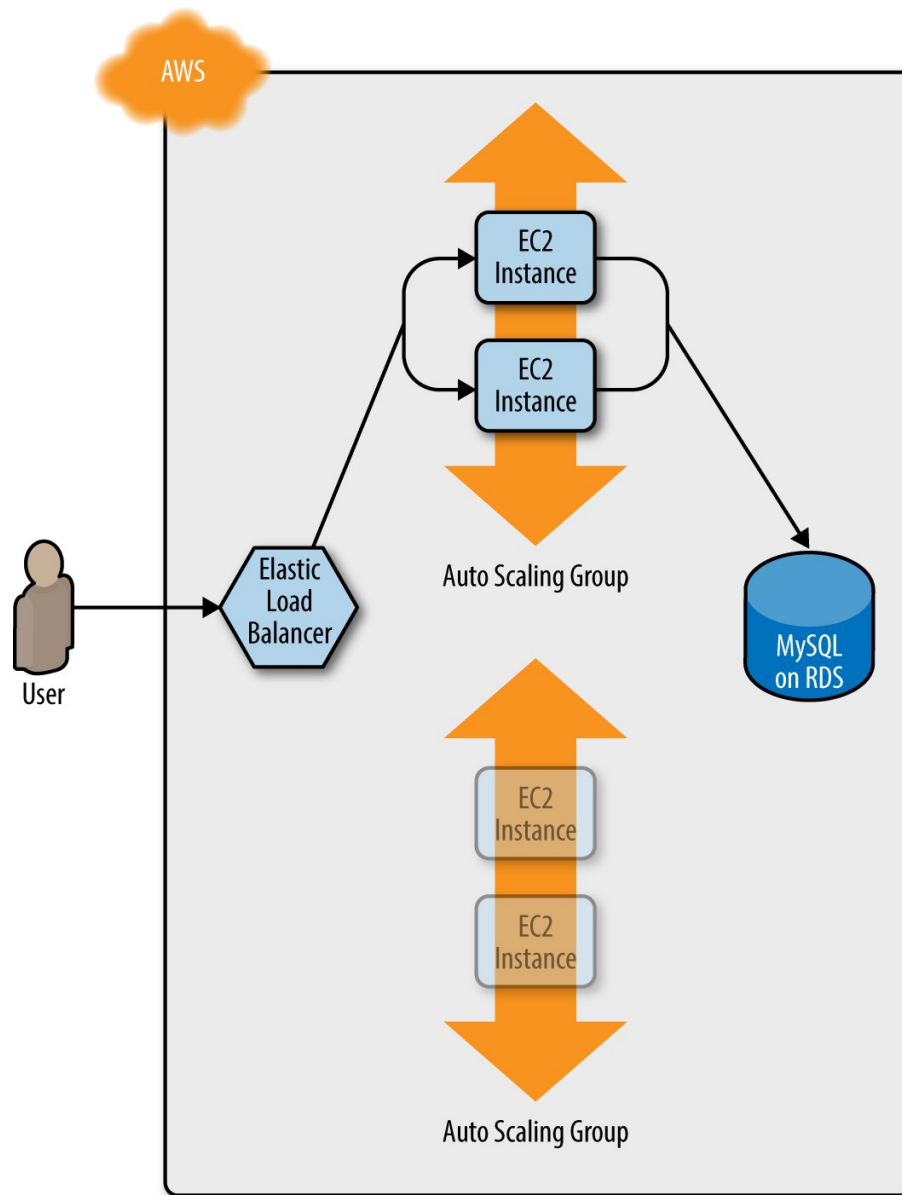
Zero-Downtime Deployment

- ◆ The way to accomplish that is to create the replacement ASG first and then destroy the original one
- ◆ Configure the name parameter of the ASG to depend directly on the name of the launch configuration *Each time the launch configuration changes (which it will when you update the AMI or User Data), its name changes, and therefore the ASG's name will change, which forces Terraform to replace the ASG
- ◆ Set the `create_before_destroy` parameter of the ASG to true, so that each time Terraform tries to replace it, it will create the replacement ASG before destroying the original
- ◆ Set the `min_elb_capacity` parameter of the ASG to the `min_size` of the cluster so that Terraform will wait for at least that many servers from the new ASG to pass health checks in the ALB before it will begin destroying the original ASG

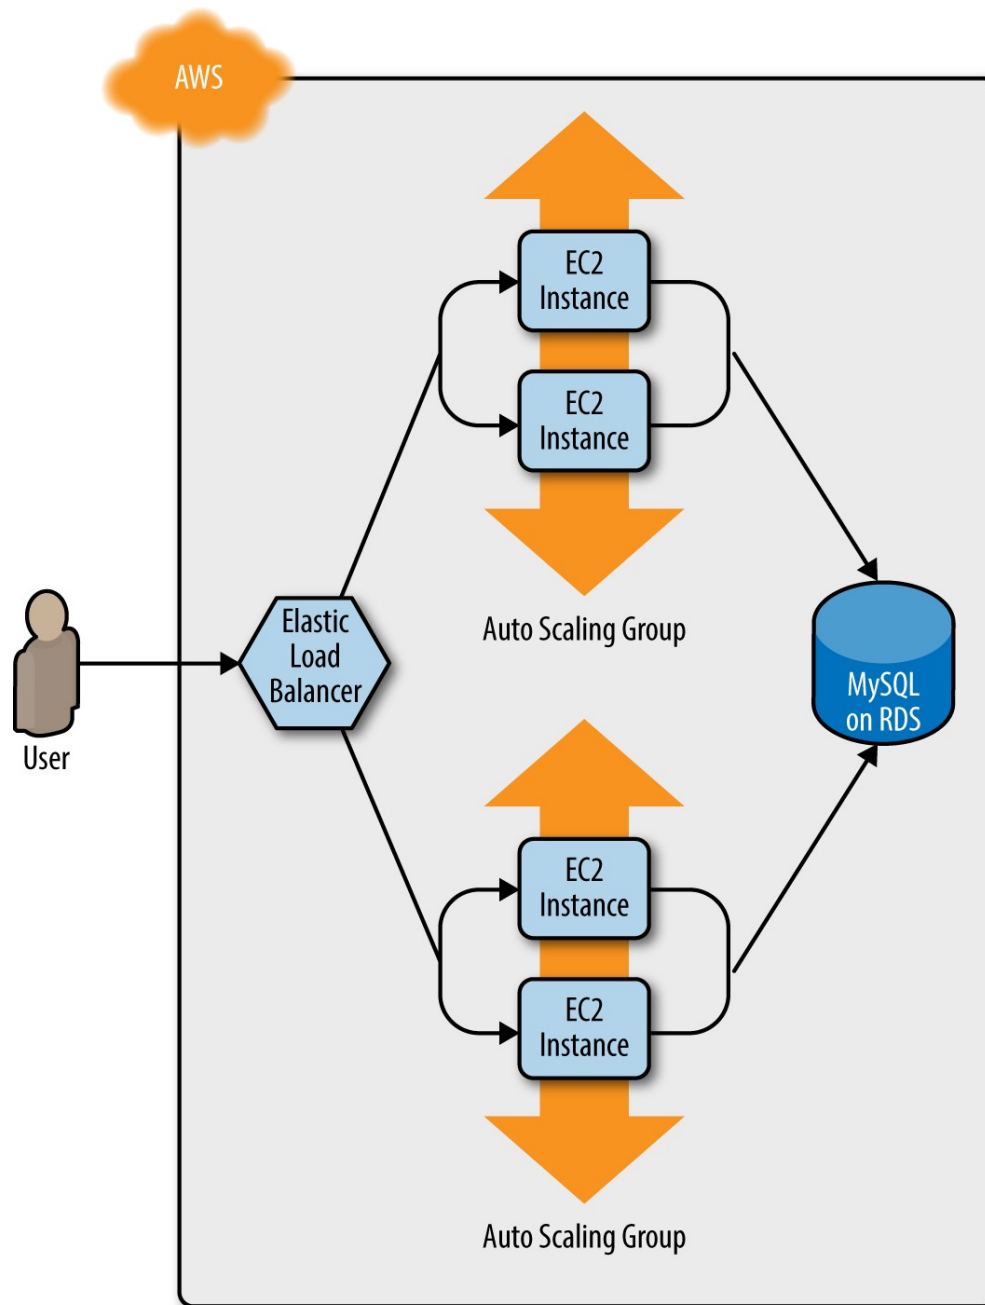
Zero-Downtime Deployment 1



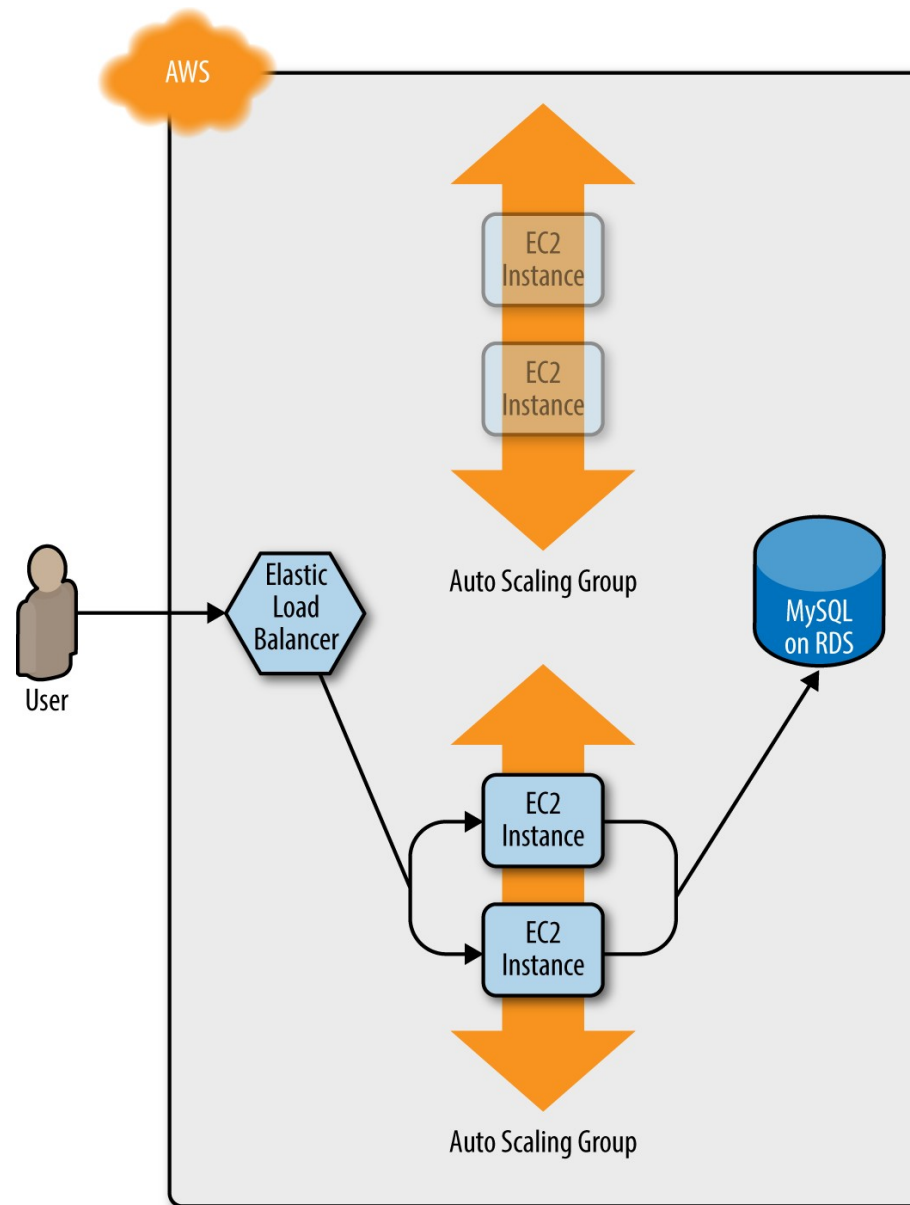
Zero-Downtime Deployment 2



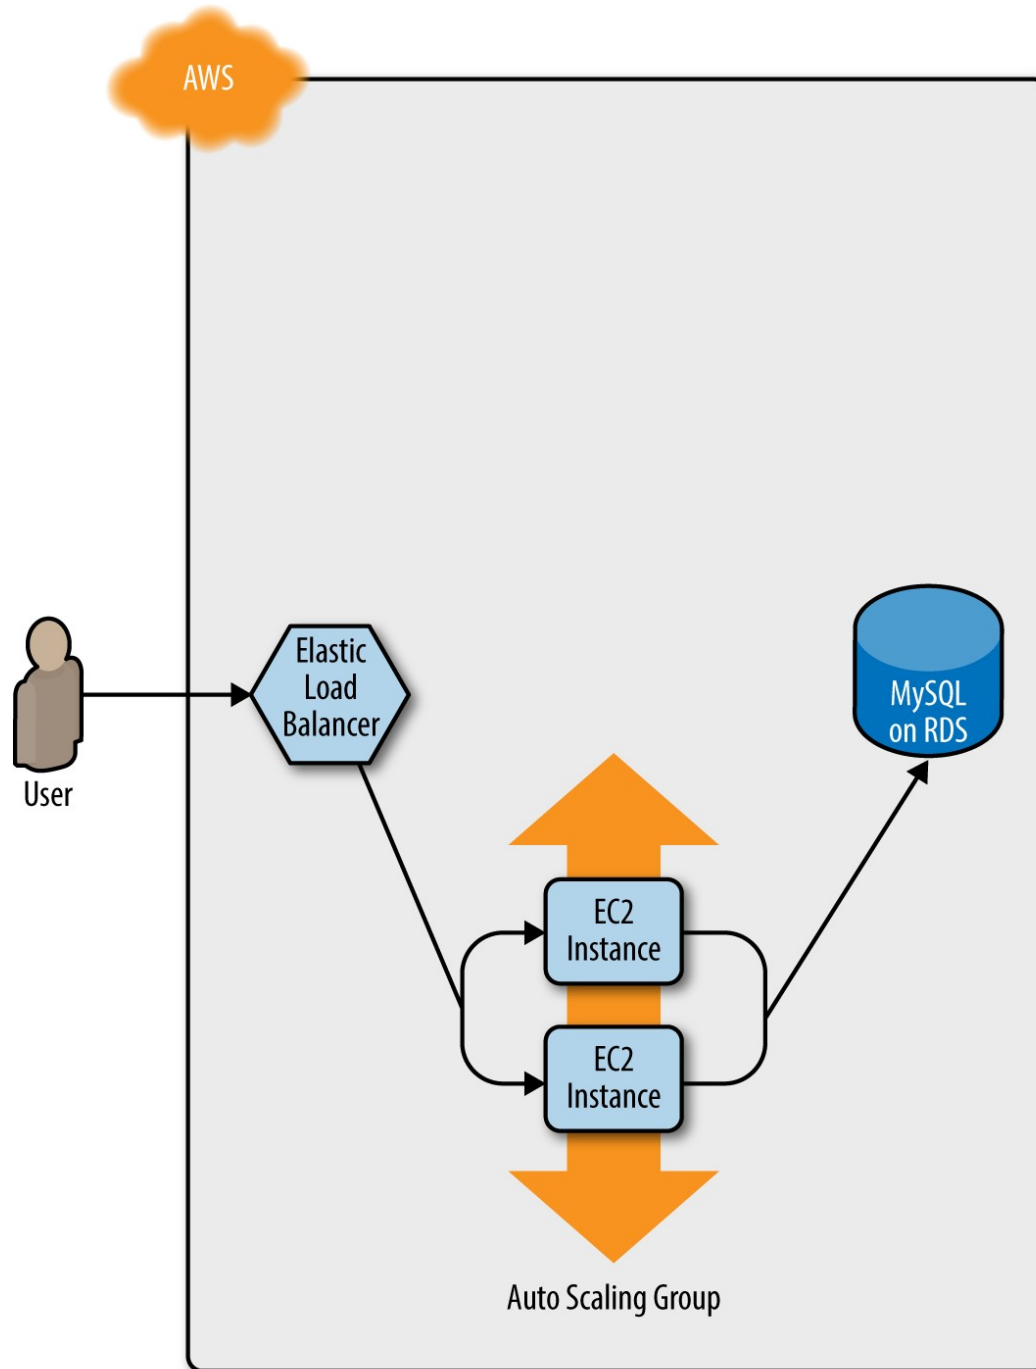
Zero-Downtime Deployment 3



Zero-Downtime Deployment 4



Zero-Downtime Deployment 5



Terraform Gotchas

- ◆ We now take a step back and point out a few gotchas
- ◆ count and for_each have limitations
 - You cannot reference any resource outputs in count or for_each
 - You cannot use count or for_each within a module configuration
- ◆ Zero-downtime deployment has limitations
 - it doesn't work with auto scaling policies
 - it resets your ASG size back to its min_size after each deployment
- ◆ Valid plans can fail
 - Terraform only looks at resources in the state file and doesn't take into account other resources
 - Plans that look good may fail because of resource conflicts
 - Ideally infrastructure should only rely on Terraform
 - Import existing infrastructure

Terraform Gotchas

- ◆ Refactoring can be tricky
 - Changes can have major effects
 - Changing the name parameter of certain resources will delete the old version of the resource and create a new version to replace it (immutable infrastructure)
- ◆ Refactoring points:
 - Always use the plan command
 - Create before destroy
 - Keep in mind that changing identifiers requires changing state
 - Some parameters are immutable so changing them requires replacing the resource

Terraform Gotchas

- ◆ Eventual consistency is consistent... eventually
- ◆ APIs for some cloud providers, such as AWS, are asynchronous and eventually consistent
 - Asynchronous means that the API might send a response immediately, without waiting for the requested action to complete
 - Eventually consistent means that it takes time for a change to propagate throughout the entire system
 - For some period of time, you might get inconsistent responses depending on which data store replica happens to respond to your API calls
- ◆ Generally, re-running `terraform apply` solves the problem