


# How to Test Terraform Code

# DevOps World is Full of Fear



# The Plan

- ◆ Manual tests
  - Manual testing basics
  - Cleaning up after tests
- ◆ Automated tests
  - Unit tests
  - Integration tests
  - End-to-end tests
  - Other testing approaches

# Manual Tests

- ◆ What does manual testing mean in Terraform?
  - How much is it like manual testing in a programming language?
- ◆ When testing Terraform code, there is no `localhost`
  - True of most IaC tools
  - You have to deploy to a real environment
  - But *not* a production environment
- ◆ This is why it is essential to have examples which can be manually tested in an real environment

# Validation Clients

- ◆ When we run tests, we normally are using some appropriate client to validate the result of the test
  - For web apps, we use 'curl' or some other tool to check the output at a specific address or URL
  - For databases, we would use a client to run queries on the db so that we can validate the results
- ◆ Best practice is to set up a test sandbox
  - There will be a lot a building and tearing down of code
  - Each developer should have their own sandbox
  - The Gold standard would be separate AWS accounts

# Cleaning Up After Testing

- ◆ Regularly clean up your sandbox environments
  - Running deployments cost money
  - It's easy to overlook infrastructure so that it sort of just hangs around
- ◆ At a minimum, use `terraform destroy` before stopping your tests
  - Also consider a regular "scrubbing" of the workspace using a cron job
- ◆ Use tools to help line
  - **cloud-nuke** : An open source tool that can delete all the resources in your cloud environment
  - **Janitor Monkey** : An open source tool that cleans up AWS resources on a configurable schedule
  - **aws-nuke** : An open source tool dedicated to deleting everything in an AWS account

# Automated Testing

- ◆ There are three kinds of automated tests:
- ◆ Unit tests
  - Unit tests verify the functionality of a single, small unit of code
  - External dependencies are replaced with test mocks
- ◆ Integration tests
  - Integration tests verify that multiple units work together correctly
  - We generally mock out other parts of the system we are not testing
- ◆ End to end testing
  - End-to-end tests involve running your entire architecture from the end-user's perspective
  - Typically use real systems everywhere, without any mocks, in an architecture that mirrors production

# Unit Testing

- ◆ The first step is to identify what a “unit” is in the Terraform world
  - A unit would be a single generic module such as the alb module
  - We can't do pure unit testing in the sense of a programming language in Terraform
  - Most unit testing is designed to work with procedural languages, not declarative ones
- ◆ Because we are deploying the module into a real environment, we are essentially doing an integration test
  - But controlling the environment is a lot like mocking it out



# Unit Testing Strategy

- ◆ This means that the basic strategy for writing unit tests for Terraform is:
  - Create a generic, standalone module
  - Create an easy-to-deploy example for that module
  - Run terraform apply to deploy the example into a real environment
  - Validate that what you just deployed works as expected
  - Run terraform destroy at the end of the test to clean up.
- ◆ For automated testing, first write the manual test, then automate
  - Remember to try the automated tests on code that has no errors to ensure the tests are being automated correctly

# TerraTest

- ◆ This is an automated test tool written in Go (the same language Terraform is written in)
- ◆ Installation steps
  - Install Go: <https://golang.org/doc/install>
  - Configure the GOPATH environment variable: <https://golang.org/doc/code.html#GOPATH>
  - Add \$GOPATH/bin to your PATH environment variable.
  - Install Dep, a dependency manager for Go: <https://golang.github.io/dep/docs/installation.html>
  - Create a folder within your GOPATH for your test code: e.g., the default GOPATH is \$HOME/go, so you could create \$HOME/go/src/terraform-up-and-running
  - Run dep init in the folder you just created. This should create Gopkg.toml, Gopkg.lock, and an empty vendors folder

# Test Installation

- ◆ Use the following Go code

```
1 package test
2
3 import (
4     "fmt"
5     "testing"
6 )
7
8 func TestGoIsWorking(t *testing.T) {
9     fmt.Println()
10    fmt.Println("If you see this text, it's working!")
11    fmt.Println()
12 }
```

- ◆ Run this test using the `go test` command and make sure you see the following output:

```
1 $ go test -v
2
3 If you see this text, it's working!
4
5 PASS
6 ok      terraform-up-and-running    0.004s
```

# Creating a

- ◆ Set up the test skeleton and set the options parameter to the alb sample directory

```
1 package test
2
3 import (
4     "github.com/gruntwork-io/terratest/modules/terraform"
5     "testing"
6 )
7
8 func TestAlbExample(t *testing.T) {
9     opts := &terraform.Options{
10         // You should update this relative path to point at your alb
11         // example directory!
12         TerraformDir: "../examples/alb",
13     }
14 }
```

```
1 $ dep ensure -add github.com/gruntwork-io/terratest/modules/terraform@v0.15.9
```

# Add code for

- ◆ We add the code to run the Terraform 'init' and 'apply' commands

```
1  func TestAlbExample(t *testing.T) {  
2      opts := &terraform.Options{  
3          // You should update this relative path to point at your alb  
4          // example directory!  
5          TerraformDir: "../examples/alb",  
6      }  
7  
8      terraform.Init(t, opts)  
9      terraform.Apply(t, opts)  
10     // or  
11     // terraform.InitAndApply(t, opts)  
12 }
```

# Testing Output Variables

- ◆ There is a testable output from here

```
1 output "alb_dns_name" {
2     value      = module.alb.alb_dns_name
3     description = "The domain name of the load balancer"
4 }
```

```
1 func TestAlbExample(t *testing.T) {
2     opts := &terraform.Options{
3         // You should update this relative path to point at your alb
4         // example directory!
5         TerraformDir: "../examples/alb",
6     }
7
8     // Deploy the example
9     terraform.InitAndApply(t, opts)
10
11     // Get the URL of the ALB
12     albDnsName := terraform.OutputRequired(t, opts, "alb_dns_name")
13     url := fmt.Sprintf("http://%s", albDnsName)
14 }
```

# The Expected Values

- ◆ A critical part of every test is the expected values that we use to compare our actual results to
- ◆ We can add those in the test code along with a method that will make an HTTP request

```
1 func TestAlbExample(t *testing.T) {  
2     (...)  
3     terraform.InitAndApply(t, opts)  
4  
5     // Get the URL of the ALB  
6     albDnsName := terraform.OutputRequired(t, opts, "alb_dns_name")  
7     url := fmt.Sprintf("http://%s", albDnsName)  
8  
9     // Test that the ALB's default action is working and returns a 404  
10  
11     expectedStatus := 404  
12     expectedBody := "404: page not found"  
13  
14     http_helper.HttpGetWithValidation(t, url, expectedStatus, expectedBody)  
15 }
```

# Timing Issues

- ◆ Because of timing issues, our tests may fail because they are run before the app is fully ready
  - We can set some parameters to retry the request a number of times

```
1 func TestAlbExample(t *testing.T) {  
2     (...)  
3     expectedStatus := 404  
4     expectedBody := "404: page not found"  
5  
6     maxRetries := 10  
7     timeBetweenRetries := 10 * time.Second  
8  
9     http_helper.HttpGetWithRetry(  
10         t,  
11         url,  
12         expectedStatus,  
13         expectedBody,  
14         maxRetries,  
15         timeBetweenRetries,  
16     )  
17 }
```



# The Test Framework

- ◆ Go has a built-in test system which is keeping track of the test being run (the `*testing.T` struct) and will report the results
- ◆ We also have to clean up and call `terraform destroy`

```
1 func TestAlbExample(t *testing.T) {  
2     (...)  
3  
4     // Clean up everything at the end of the test  
5     defer terraform.Destroy(t, opts)  
6  
7     // Deploy the example  
8     terraform.InitAndApply(t, opts)  
9  
10    // Get the URL of the ALB  
11    (...)  
12  
13 }
```

# Automated Testing Pointers

- ◆ Manual testing should be done in a sandbox account
  - For automated testing, this is even more important
  - A totally separate account is recommended.
  - As your automated test suite grows, you might be spinning up hundreds or thousands of resources in every test suite, so keeping them isolated from everything else is essential
  - Teams should consider a completely separate environment just for automated testing
  - This is separate even from the sandbox environments you use for manual testing.
- ◆ Go test can be run with a time limit (10 minutes by default) which it kills the the test environment
  - At this point all the tests fail but the environment is not destroyed

```
1 | go test -v -timeout 30m
```

# Recap - Key Parts of the Test

- ◆ Running terraform init
- ◆ Running terraform apply
- ◆ Reading output variables using terraform output
- ◆ Repeatedly making HTTP requests to the ALB
- ◆ Running terraform destroy

# Integration Tests

- ◆ We have been treating modules as our basic units in Terraform
- ◆ An integration test would deploy several modules and see that they work correctly
  - If the modules should not have anything in them hard-coded for the staging environment
- ◆ We never do any integration tests unless all of the modules have been thoroughly unit tested
- ◆ To do an integration test of the webserver and the database backend:
  - First deploy the `mysql` server
  - Then deploy the `hello-world-app`
  - Run the test
  - Undeploy the `hello-world-app`
  - Undeploy `mysql`

# Running the Integration Test

- ◆ The test code to run the integration test would look something like this:

```
1 // Replace these with the proper paths to your modules
2 const dbDirStage = "../live/stage/data-stores/mysql"
3 const appDirStage = "../live/stage/services/hello-world-app"
4
5 func TestHelloWorldAppStage(t *testing.T) {
6     t.Parallel()
7
8     // Deploy the MySQL DB
9     dbOpts := createDbOpts(t, dbDirStage)
10    defer terraform.Destroy(t, dbOpts)
11    terraform.InitAndApply(t, dbOpts)
12
13    // Deploy the hello-world-app
14    helloOpts := createHelloOpts(dbOpts, appDirStage)
15    defer terraform.Destroy(t, helloOpts)
16    terraform.InitAndApply(t, helloOpts)
17
18    // Validate the hello-world-app works
19    validateHelloApp(t, helloOpts)
20 }
```

- ◆ One issue that has to be dealt with is where each component is storing their state
  - We don't want the testing to overwrite the actual state file

# Faster Integration Testing

- ◆ Running the stages mentioned earlier may be unnecessary if all we are doing is making changes to the `hello-world-app`
- ◆ In this case we can eliminate overhead by doing:
  - Run `terraform apply` on the `mysql` module
  - Run `terraform apply` on the `hello-world-app` module
  - Work on the module
    - Make changes to `hello-world-app`
    - Run `terraform apply` on the `hello-world-app` module to implement updates
    - Validate
    - If it all works go to step 4, else go back to step 3
  - Run `terraform destroy` on the `hello-world-app` module
  - Run `terraform destroy` on the `mysql` module

# Terratest Support

- ◆ Terratest supports this natively with the `test_structure` package
  - Each stage of your test in a function with a name
  - Terratest can skip some of those names by setting environment variables
  - Each test stage stores test data on disk so that it can be read back from disk on subsequent test runs

```
1 func TestHelloWorldAppStageWithStages(t *testing.T) {  
2     t.Parallel()  
3  
4     // Store the function in a short variable name solely to make the  
5     // code examples fit better in the book.  
6     stage := test_structure.RunTestStage  
7  
8     // Deploy the MySQL DB  
9     defer stage(t, "teardown_db", func() { teardownDb(t, dbDirStage) })  
10    stage(t, "deploy_db", func() { deployDb(t, dbDirStage) })  
11  
12    // Deploy the hello-world-app  
13    defer stage(t, "teardown_app", func() { teardownApp(t, appDirStage) })  
14    stage(t, "deploy_app", func() { deployApp(t, dbDirStage, appDirStage) })  
15  
16    // Validate the hello-world-app works  
17    stage(t, "validate_app", func() { validateApp(t, appDirStage) })  
18 }
```

# Test Stages

- ◆ The `RunTestStage` method takes three arguments:
  - `t` : the built in `test` structure that manages the state of the test
    - for example, calling `t.Fail()` causes the test to fail
  - `Stage name` : the name for this test stage
  - `Code to execute` : Any function to execute
- ◆ The code to implement `deployDB` and `teardownDb` would be

```
1 func deployDb(t *testing.T, dbDir string) {  
2     dbOpts := createDbOpts(t, dbDir)  
3  
4     // Save data to disk so that other test stages executed at a later  
5     // time can read the data back in  
6     test_structure.SaveTerraformOptions(t, dbDir, dbOpts)  
7  
8     terraform.InitAndApply(t, dbOpts)  
9 }
```

```
1 func teardownDb(t *testing.T, dbDir string) {  
2     dbOpts := test_structure.LoadTerraformOptions(t, dbDir)  
3     defer terraform.Destroy(t, dbOpts)  
4 }
```



# Selective Test Stages

- ◆ Staging allows for selective skipping of stages

```
1 $ SKIP_teardown_db=true \  
2   SKIP_teardown_app=true \  
3   go test -timeout 30m -run 'TestHelloWorldAppStageWithStages' \  
4   \  
5   (...) \  
6   \  
7   The 'SKIP_deploy_db' environment variable is not set, \  
8   so executing stage 'deploy_db'. \  
9   \  
10  (...) \  
11  The 'teardown_app' environment variable is set, \  
12  so skipping stage 'deploy_db'.
```

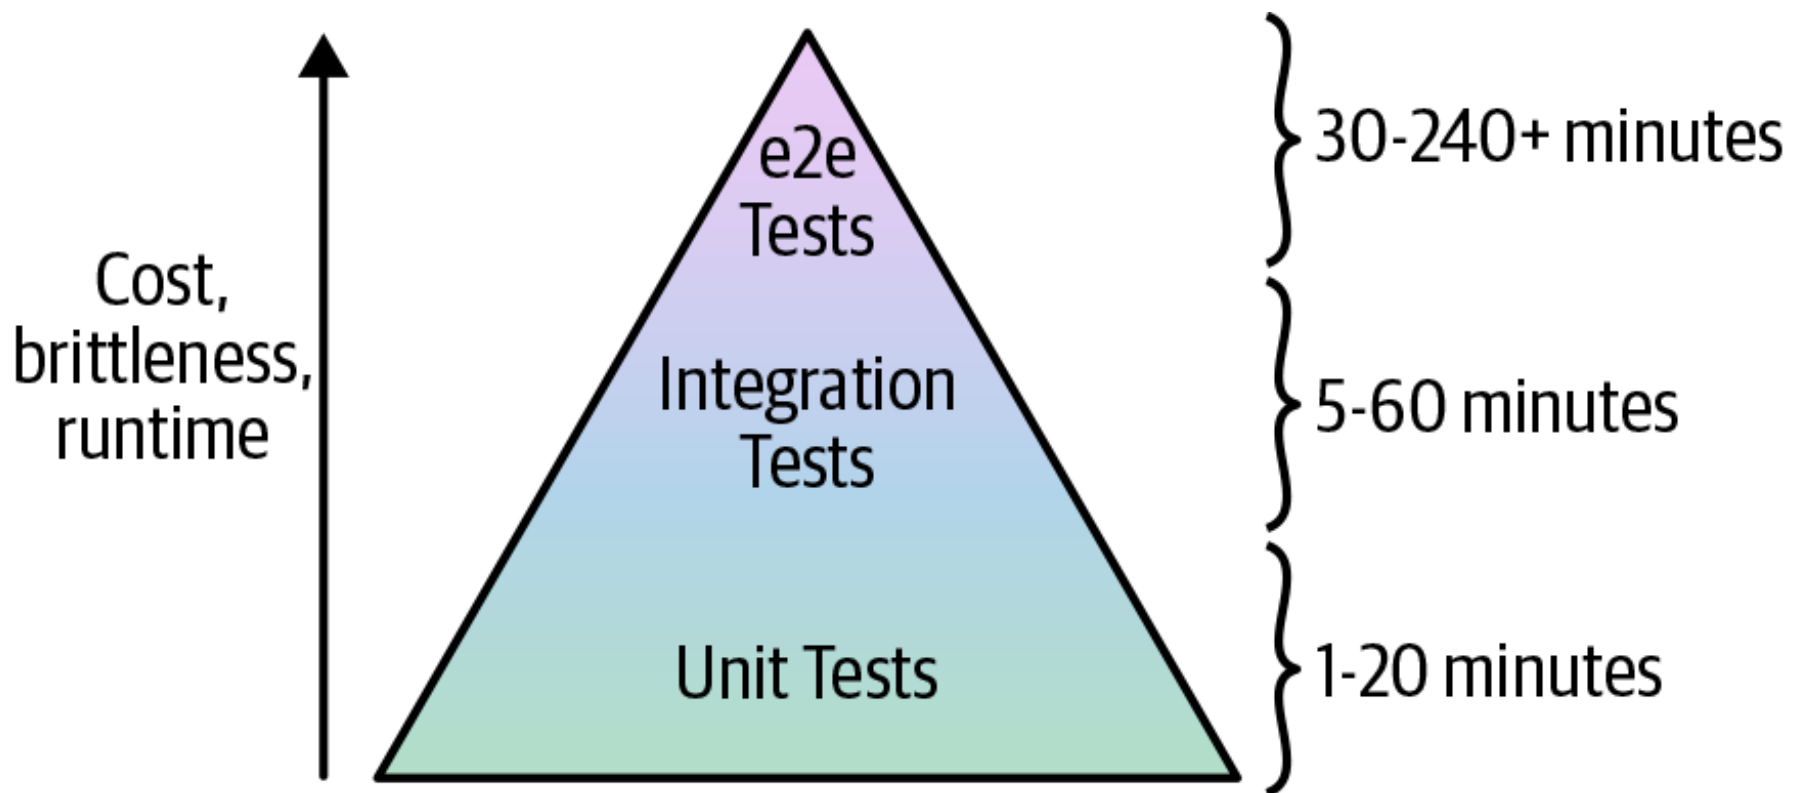
# Retries

- ◆ Tests can fail for various transient reasons that are random and unpredictable
  - A common way of dealing with false positives is retrying the test on a failure
- ◆ Terratest supports automatic retries by setting "retry" options

```
1 func createHelloOpts(  
2 dbOpts *terraform.Options,  
3 terraformDir string) *terraform.Options {  
4  
5     return &terraform.Options{  
6         TerraformDir: terraformDir,  
7  
8         Vars: map[string]interface{}{  
9             "db_remote_state_bucket": dbOpts.BackendConfig["bucket"],  
10            "db_remote_state_key": dbOpts.BackendConfig["key"],  
11            "environment": dbOpts.Vars["db_name"],  
12        },  
13  
14        // Retry up to 3 times, with 5 seconds between retries,  
15        // on known errors  
16        MaxRetries: 3,  
17        TimeBetweenRetries: 5 * time.Second,  
18        RetryableTerraformErrors: map[string]string{  
19            "RequestError: send request failed": "Throttling issue?",  
20        },  
21    }  
22 }
```

# End to End Tests

- ◆ As our tests include more and more code, they become longer to execute and more costly to set up and run
- ◆ We should plan for a large number of unit tests, smaller number of integration tests and even smaller number of end-to-end tests



# End to End Setup

- ◆ With larger and more complicated infrastructure, setting up a stable infrastructure (test environments, namespaces etc) this only becomes more difficult
  - Therefore, you want to do as much of your testing as low in the pyramid as you can because the bottom of the pyramid offers the fastest, most reliable feedback loop
- ◆ Deploying a complicated architecture from scratch is untenable for several reasons
- ◆ Too slow:
  - The more complex the infrastructure, longer it takes to set up
  - Limits the amount of testing that can be done which means slow feedback
- ◆ Too brittle:
  - Constantly redeploying a complex setup increases the likelihood of transient errors
  - This means constant retries which inhibit the whole testing effort

# End to End Strategy

- ◆ A common end-to-end strategy is:
  - A persistent, production-like environment called “test” is deployed which is left running
  - Every time a change is made to the infrastructure, the end-to-end test does the following:
    - Applies the infrastructure change to the test environment
    - Runs validations against the test environment (e.g., use Selenium to test your code from the end-user’s perspective) to make sure everything is working
- ◆ More closely mimics how you’ll be deploying those changes in production
  - Also confirms the deployment process also works - for example, the change can be made with zero downtime

# Static Analysis

- ◆ Static analysis involves running tools that examine the structure of the code without executing it
- ◆ Common tools are:
  - `terraform validate` : a command built into Terraform that you can use to check your Terraform syntax and types
  - `tflint` : A “lint” tool for Terraform that can scan Terraform code and catch common errors and potential bugs based on a set of built-in rules
  - HashiCorp Sentinel : A “policy as code” framework that allows you to enforce rules across various HashiCorp tools

# Property Testing

- ◆ These are testing tools like `rspec-terraform` that use Domain Specific Languages to confirm that the infrastructure conforms to a specification
- ◆ For example:

```
1 describe file('/etc/myapp.conf') do
2     it { should exist }
3     its('mode') { should cmp 0644 }
4 end
5
6 describe apache_conf do
7     its('Listen') { should cmp 8080 }
8 end
9
10 describe port(8080) do
11     it { should be_listening }
12 end
```

# Key Takeaways

- ◆ When testing Terraform code, there is no localhost
  - All manual testing is done by deploying real resources into one or more isolated sandbox environments
- ◆ Regularly clean up your sandbox environments
  - Otherwise, the environments will become unmanageable, and costs will spiral out of control
- ◆ You cannot do pure unit testing for Terraform code
  - Therefore, all automated testing is done by writing code that deploys real resources into one or more isolated sandbox environments
- ◆ You must namespace all of your resources
  - This ensures that multiple tests running in parallel do not conflict with one another
- ◆ Smaller modules are easier and faster to test
  - Smaller modules are easier to create, maintain, use, and test.